58	A Mule application is designed to call the Google Maps API to perform a distance computation, and is deployed to CloudHub. At a minimum, what (if anything) should be configured in the TLS context of the HTTP Request configuration to meet these requirements? A. R st a private key from Google, create a PKCS# 12 file with it, and add it in KeyStore as part of the TLS context B. Do he Google public certificate from a browser, generate a JKS file from it, and add it in KeyStore as part LS context C. The con is built-in and nothing extra is required for the TLS context D. Download the blic certificate from a browser, generate a JKS file from it, and add it in TrustStore t of the TLS context	D	https://docs.mulesoft.com/mule-runtime/latest/tls-configuration
59	A project team is working on an API impleme ing an API RAML specification as a s point. The team has updated the RAML specification to include new operations, and has pu new version to Exchange. M ile, another team is working on an API client Mule application to consume the same API implementatio During this development, what must be performed by the Mule on team advantage of the newly added operations? A. Scaffold the API client application with APIkit using the new RAML B. Upd API connector in the API impleme and publish to Exc C. Update connector from Exchange API client applica D. Copy the API tation's updated component into th client	C	
60	A marketing organization is e application to r campaign data. The Mule ap will periodically check for a file in a SFTP location and process the re he file. The siz file can vary from 10MB to 5GB. Due ited availability of vCores, the Mule application is deployed gle h wo nfigured with vCore size 0.2. The application must trans d send diff s of this file to three different downstream SFTP locatio What is the most idioma ed for its intend d performant way to configure the SFTP operations or eve to process the large files to sup ese deployment ments. A. Use an in-mem n-repeatable st B. Use a file-stored non-repeatable stream C. Use an in-memory repeatable stream D. Use a file-stored repeatable stream	D	https://docs.mulesoft.com/mule-runtime/4.4/streaming-about
61	A corporation has deployed multiple Mule applications implementing various public and private APIs to different CloudHub workers. These APIs are critical applications that must be highly available and inline with the reliability SLA as defined by stakeholders. How can API availability (liveness or readiness) be monitored so the Ops team receives outage notifications? A. Enable monitoring of individual applications from Anypoint Monitoring B. Configure alerts with failure conditions in Runtime Manager C. Configure alerts with failure conditions in API Manager D. Use Anypoint Functional Monitoring to test APIs' functional behavior	D	https://docs.mulesoft.com/monitoring/ https://docs.mulesoft.com/api-functional-monitoring/afm-in-anypoint-platform
62	A Mule application must periodically process a large data set, which varies from 6GB-8GB, from a back-end database and write transformed data to an FTPS server using a property configured batch job scope. The performance requirements of the application are approved to run in a CloudHub 0.2 vCore with 8GB storage capacity, and the concurrency requirements are met. How can the high rate of records be effectively managed this application? A. Use streaming with a file store repeatable strategy for reading records from the database, and batch aggregator with streaming to write to FTPS B. Use streaming with an in-memory repeatable store strategy for reading records from the database, and batch aggregator with streaming to write to FTPS C. Use streaming with a file store repeatable strategy for reading records from the database and batch aggregator with an optimal size D. Use streaming with a file store repeatable strategy for reading records from the database and batch aggregator without any required configuration	A	
63	ACME Tractors is designing a Mule application named Inventory that uses a persistent object store. The Inventory Mule application is deployed to CloudHub and is configured to use Object Store v2. Another Mule application named Cleanup is being developed to delete values from the Inventory Mule application's persistent object store. The Cleanup Mule application will also be deployed to CloudHub. What is the most direct way for the Cleanup Mule application to delete values from the Inventory Mule application's persistent object store with the least latency? A. Use a VM connector configured to directly access the persistent queue of the Inventory Mule application's persistent object store B. Use an Object Store connector configured to access the Inventory Mule application's persistent object store C. Use the Object Store v2 REST API configured to access the Inventory Mule application's persistent object store D. Use an Anypoint MQ connector configured to directly access the Inventory Mule application's persistent object store	C	https://docs.mulesoft.com/object-store/osv2-guide
64	A Mule application is running on a customer-hosted Mule runtime in an organization's network. The Mule application acts as a producer of asynchronous Mule events. Each Mule event must be broadcast to all interested external consumers outside the Mule application. The Mule events should be published in a way that is guaranteed in normal situations and also minimizes duplicate delivery in less frequent failure scenarios. The organizational firewall is configured to only allow outbound traffic on ports 80 and 443. Some external event consumers are within the organizational network, while others are located outside the firewall. What Anypoint Platform service is most idiomatic (used for its intended purpose) for publishing these Mule events to all external consumers while addressing the desired reliability goals? A. CloudHub VM queues B. CloudHub Shared Load Balancer C. Anypoint MQ D. Anypoint Exchange	C	https://docs.mulesoft.com/mq/

65	A Mule application is built to support a local transaction for a **series of operations on a single database**. The Mule application has a Scatter-Gather that participates in the local transaction. What is the behavior of the Scatter-Gather when running within this local transaction? A. Execution of all routes within the Scatter-Gather occurs in parallel Any error that occurs inside the Scatter-Gather will result in a rollback of all the database operations B. Execution of each route within the Scatter-Gather occurs sequentially Any error that occurs inside the Scatter-Gather will NOT result in a rollback of any of the database operations C. Execution of each route within the Scatter-Gather occurs in parallel Any error that occurs inside the Scatter-Gather will NOT result in a rollback of any of the database operations D. Execution of each route within the Scatter-Gather occurs sequentially Any error that occurs inside the Scatter-Gather will result in a rollback of all the database operations	D	
66	A Mule application is being designed for deployment to a single CloudHub worker. The Mule application will have a flow that connects to a SaaS system to perform some operations each time the flow is invoked. The SaaS system connector has operations that can be configured to request a short-lived token (fifteen minutes) that can be reused for subsequent connections within the fifteen minute time window. After the token expires, a new token must be requested and stored. What is the most performant and idiomatic (used for its intended purpose) Anypoint Platform component or service to use to support persisting and reusing tokens in the Mule application to help speed up reconnecting the Mule application to the SaaS application? A. Persistent object store B. Variable C. Database D. Non persistent object store	D	
67	To implement predictive maintenance on its machinery equipment, ACME Tractors has installed thousands of IoT sensors that will send data for each machinery asset as sequences of JMS messages, in near real-time, to a JMS queue named SENSOR_DATA on a JMS server. The Mule application contains a JMS Listener operation configured to receive incoming messages from the JMS server's SENSOR_DATA JMS queue. The Mule application persists each received JMS message, then sends a transformed version of the corresponding Mule event to the machinery equipment back-end systems. The Mule application will be deployed to a multi-node, customer-hosted Mule runtime cluster. Under normal conditions, each JMS message should be processed exactly once. How should the JMS Listener be configured to maximize performance and concurrent message processing of the JMS queue? A. Set numberOfConsumers to a value greater than one Set primaryNodeOnly = false B. Set numberOfConsumers = 1 Set primaryNodeOnly = true C. Set numberOfConsumers = 1 Set primaryNodeOnly = false D. Set numberOfConsumers to a value greater than one Set primaryNodeOnly = true	C	
68	The AnyAirline organization's passenger reservations center is designing an integration solution that combines invocations of three different System APIs (bookFlight, bookHotel, and bookCar) in a business transaction. Each System API makes calls to a single database. The entire business transaction **must be rolled back when at least one of the APIs fails**. What is the most idiomatic (used for its intended purpose) way to integrate these APIs in near real-time that provides the best balance of consistency, performance, and reliability? A. Implement local transactions within each API implementation Configure each API implementation to also participate in the same extended Architecture (XA) transaction Implement caching in each API implementation to improve performance B. Implement local transactions in each API implementation Coordinate between the API implementations using a Saga pattern Apply various compensating actions depending on where a failure occurs C. Implement extended Architecture (XA) transactions between the API implementations Coordinate between the API implementations using a Saga pattern Implement caching in each API implementation to improve performance D. Implement an extended Architecture (XA) transaction manager in a Mule application using a Saga pattern Connect each API implementation with the Mule application using XA transactions Apply various compensating actions depending on where a failure occurs	B	
69	ACME Tractors is creating a Mule application that will be deployed to CloudHub. The Mule application has a property named dbPassword that stores a database user's password. The organization's security standards indicate that the dbPassword property must be hidden from every Anypoint Platform user after the value is set in the Runtime Manager Properties tab. What configuration in the Mule application helps hide the dbPassword property value in Runtime Manager? A. Add the dbpassword property to the secureProperties section of the mule-artifact.json file B. Use secure::dbPassword as the property placeholder name and store the cleartext (unencrypted) value in a secure properties placeholder file C. Use secure::dbPassword as the property placeholder name and store the properly encrypted value in a secure properties placeholder file D. Add the dbpassword property to the secureProperties section of the pom.xml file	A	https://docs.mulesoft.com/cloudhub/secure-application-properties#create-safely-hidden-application-properties

70	An organization has previously provisioned its own AWS virtual private cloud (VPC) that contains several AWS instances. The organization now needs to use CloudHub to host a Mule application that will implement a REST API. Once deployed to CloudHub, this Mule application must be able to communicate securely with the customer-provisioned AWS VPC resources within the same region, without being interceptable on the public internet. What Anypoint Platform features should be used to meet these network communication requirements between CloudHub and the existing customer-provisioned AWS VPC? A. Configure a MuleSoft-hosted (CloudHub) dedicated load balancer with mapping rules that allow secure traffic from the range of IP addresses located in the customer-provisioned AWS VPC to access the Mule application B. Add a MuleSoft-hosted (CloudHub) Anypoint VPC configured with VPC peering to the range of IP addresses located in the customer-provisioned AWS VPC C. Add default API allowlist policies to API Manager to automatically secure traffic from the range of IP addresses located in the customer-provisioned AWS VPC to access the Mule application D. Configure an external identity provider (IdP) in Anypoint Platform with certificates from an AWS Transit Gateway for the customer-hosted AWS VPC, where the certificates allow the range of IP addresses located in the customer- provisioned AWS VPC	B	
71	An API is being implemented using the components of Anypoint Platform. The API implementation needs to be managed and governed (by applying API policies) on Anypoint Platform. What needs to be done in Anypoint Platform before the API implementation can be governed by Anypoint Platform? A. The API must be shared with the potential developers through an API portal so API consumers can interact with the API B. The API must be published to Anypoint Exchange and a corresponding API instance ID must be obtained from API Manager to be used in the API implementation C. The API implementation source code must be committed to a source control management system (such as GitHub) D. A RAML definition of the API must be created in API Designer so it can then be published to Anypoint Exchange	B	
72	ACME Tractors has implemented a continuous integration (CI) lifecycle that promotes Mule applications through code, build, and 4 test stages. To standardize the organization's CI journey, a new dependency control approach is being designed to store artifacts that include information such as dependencies, versioning, and build promotions. To implement these process improvements, the organization will now require developers to maintain all dependencies related to Mule application code in a shared location. What is the most idiomatic (used for its intended purpose) type of system the organization should use in a shared location to standardize all dependencies related to Mule application code? A. The Anypoint Object Store service at cloudhub.io B. A binary artifact repository C. A MuleSoft-managed repository at repository.mulesoft.org D. API Community Manager	B	
73	ACME Tractors is designing a Mule application to periodically poll an SFTP location for new files containing sales order records and then process those sales orders. Each sales order must be processed exactly once. To support this requirement, the Mule application must identify and filter duplicate sales orders on the basis of a unique ID contained in each sales order record and then only send the new sales orders to the downstream system. What is the most idiomatic (used for its intended purpose) Anypoint connector, validator, or scope that can be configured in the Mule application to filter duplicate sales orders on the basis of the unique ID field contained in each sales order record? A. Configure a Cache scope to filter and store each record from the received file by the order ID B. Configure an Idempotent Message Validator component to filter each record by the order ID C. Configure a Database connector to filter and store each record by the order ID D. Configure a watermark in an On New or Updated File event source to filter unique records by the order ID	B	https://docs.mulesoft.com/mule-runtime/4.4/idempotent-message-validator
74	The ACME Tractors organization has deployed both Mule and non-Mule API implementations to integrate its customer and order management systems. All the APIs are available to REST clients on the public internet. The organization wants to monitor these APIs by running health checks; for example, to determine if an API can properly accept and process requests. The organization does not have subscriptions to any external monitoring tools and also does not want to extend its IT footprint. What Anypoint Platform feature provides the most idiomatic (used for its intended purpose) way to monitor the availability of both the Mule and the non-Mule API implementations? A. Runtime Manager B. API Manager C. API Functional Monitoring D. Anypoint Visualizer	C	
75	A REST API is being designed. It will be implemented as a Mule application and managed on Anypoint Platform. What **standard interface definition languag**e can be used to define REST APIs? A. Web Services Description Language (WSDL) B. AsyncAPI Specification (AsyncAPI) C. OpenAPI Specification (OAS) D. YAML Ain't Markup Language (YAML)	C	https://docs.mulesoft.com/release-notes/platform/oas3

76	An external web UI application currently accepts occasional HTTP requests from client web browsers to change (insert, update, or delete) inventory pricing information in an inventory system's database. Each inventory pricing change must be transformed and then synchronized with multiple customer experience systems in near real-time (in under 10 seconds). New customer experience systems are expected to be added in the future. The database is used heavily and limits the number of SELECT queries that can be made to the database to 10 requests per hour per user. What is the most scalable, idiomatic (used for its intended purpose), decoupled, reusable, and maintainable integration mechanism available to synchronize each inventory pricing change with the various customer experience systems in near real-time? A. Write a Mule application with a Database On Table Row event source configured for the inventory pricing database, with the watermark attribute set to an appropriate database column In the same flow, use a Scatter-Gather to call each customer experience system's REST API with transformed inventory-pricing records B. Replace the external web UI application with a Mule application to accept HTTP requests from client web browsers In the same Mule application, use a Batch Job scope to test if the database request will succeed, aggregate pricing changes within a short time window, and then update both the inventory pricing database and each customer experience system using a Parallel For Each scope C. Add a trigger to the inventory-pricing database table so that for each change to the inventory pricing database, a stored procedure is called that makes a REST call to a Mule application Write the Mule application to publish each Mule event as a message to an Anypoint MQ exchange Write other Mule applications to subscribe to the Anypoint MQ exchange, transform each received message, and then update the Mule application's corresponding customer experience system(s) D. Write a Mule application with a Database On Table Row event source configured for the inventory pricing database, with the ID attribute set to an appropriate database column In the same flow, use a Batch Job scope to publish transformed inventory-pricing records to an Anypoint MQ queue Write other Mule applications to subscribe to the Anypoint MQ queue, transform each received message, and then update the Mule application's corresponding customer experience system(s)	A	
77	A Mule application uses APIkit for SOAP to implement a SOAP web service. The Mule application has been deployed to a CloudHub worker in a testing environment. The integration testing team wants to use a SOAP client to perform integration testing. To carry out the integration tests, the integration team must obtain the interface definition for the SOAP web service. What is the most idiomatic (used for its intended purpose) way for the integration testing team to obtain the interface definition for the deployed SOAP web service in order to perform integration testing with the SOAP client? A. Retrieve the XML file(s) from Runtime Manager B. Retrieve the RAML file(s) from the deployed Mule application C. Retrieve the OpenAPI Specification file(s) from API Manager D. Retrieve the WSDL file(s) from the deployed Mule application	D	https://www.infoviewsystems.com/creating-a-soap-api-via-mulesoft/
78	An organization is designing a Mule application to support an all-or-nothing transaction between several database operations and some other connectors so that they all automatically roll back if there is a problem with any of the connectors. Besides the Database connector, what other Anypoint connector can be used in the Mule application to participate in the all-or-nothing transaction? A. JMS(or VM) B. SFTP C. Anypoint MQ D. Object Store	A	
79	ACME Tractors has defined a common object model in Java to mediate the communication between different Mule applications in a consistent way. A Mule application is being built to use this common object model to process responses from a SOAP API and a REST API and then write the processed results to an order management system. The developers want Anypoint Studio to utilize these common objects to assist in creating mappings for various transformation steps in the Mule application. What is the most idiomatic (used for its intended purpose) and performant way to utilize these common objects to map between the inbound and outbound systems in the Mule application? A. Use the WSS module B. Use the Java module C. Use JAXB (XML) and Jackson (JSON) data bindings D. Use the Transform Message component	D	
80	The ACME Best Insurance company has an Anypoint Runtime Fabric on VMs/Bare Metal (RTF-VM) appliance installed on its own customer-hosted AWS infrastructure. Mule applications are deployed to this RTF-VM appliance. As part of the company standards, the Mule application logs must be forwarded to an external log management tool (LMT). Given the company's current setup and requirements, what is the most idiomatic (used for its intended purpose) way to send Mule application logs to the external LMT? A. In RTF-VM, install and configure the external LTM's log-forwarding agent B. In RTF-VM, configure the out-of-the-box external log forwarder C. In RTF-VM, edit the pod configuration to automatically install and configure an Anypoint Monitoring agent D. In each Mule application, configure custom Log4j settings	A	https://docs.mulesoft.com/runtime-fabric/2.1/#external-log-forwarding

81	ACME Tractors has an HTTPS-enabled Mule application named Orders API that receives requests from another Mule application named Process Orders. The communication between these two Mule applications must be secured by TLS mutual authentication (two-way TLS). At a minimum, what must be stored in each truststore and keystore of these two Mule applications to properly support two-way TLS between the two Mule applications while properly protecting each Mule application's keys? A. Orders API truststore: The Process Orders public key Orders API keystore: The Orders API private key Process Orders truststore: The Orders API public key Process Orders keystore: The Process Orders private key B. Orders API truststore: The Orders API private key and public key Process Orders keystore: The Process Orders private key and public key C. Orders API truststore: The Process Orders public key Orders API keystore: The Orders API private key and public key Process Orders truststore: The Orders API public key Process Orders keystore: The Process Orders private key and public key D. Orders API truststore: The Orders API public key Process Orders keystore: The Process Orders private key and public key	C	
82	In Anypoint Platform, a company wants to configure multiple identity providers (IdPs) for multiple lines of business (LOBs). Multiple business groups, teams, and environments have been defined for these LOBS. What Anypoint Platform feature can use multiple IdPs across the company's business groups, teams, and environments? A. Client Management B. Roles and permissions C. MuleSoft-hosted (CloudHub) dedicated load balancers D. Virtual private clouds	A	https://docs.mulesoft.com/access-management/environments#add-an-idp-to-an-existing-environment
83	A Mule application is being designed to perform product orchestration. The Mule application must join together the responses from an Inventory API and a Product Sales History API with the least latency. To minimize the overall latency, what is the most idiomatic (used for its intended purpose) design to call each API request in the Mule application as quickly as possible? A. Call each API request in a separate Async scope B. Call each API request in a separate route of a Scatter-Gather C. Call each API request in a separate lookup call from a DataWeave reduce function D. Call each API request in a separate route of a Parallel For Each scope	B	
84	An ACME Farms project team is planning to build a new API that is required to work with data from different domains across the organization. The organization has a policy that all project teams should leverage existing investments by reusing existing APIs and related resources and documentation that other project teams have already developed and deployed. To support reuse, where on Anypoint Platform should the project team go to discover and read existing APIs, discover related resources and documentation, and interact with mocked versions of those APIs? A. Design Center B. API Manager C. Anypoint Exchange D. Runtime Manager	C	https://docs.mulesoft.com/design-center/design-mocking-service
85	ACME Tractors is designing an integration Mule application to process orders by submitting them to a back-end system for offline processing. Each order will be received by the Mule application through an HTTPS POST and must be acknowledged immediately. Once acknowledged, the order will be submitted to a back-end system. Orders that cannot be successfully submitted due to rejections from the back-end system will need to be processed manually (outside the back-end system). The Mule application will be deployed to a customer-hosted runtime and is able to use an existing ActiveMQ broker if needed. The ActiveMQ broker is located inside the organization's firewall. The back-end system has a track record of unreliability due to both minor network connectivity issues and longer outages. What idiomatic (used for their intended purposes) combination of Mule application components and ActiveMQ queues are required to ensure automatic submission of orders to the back-end system while supporting but minimizing manual order processing? A. One or more On Error scopes to assist calling the back-end system One or more ActiveMQ long-retry queues A persistent dead-letter object store configured in the CloudHub Object Store service B. A Batch Job scope to call the back-end system An Until Successful scope containing Object Store components for long retries A dead-letter object store configured in the Mule application C. An Until Successful scope to call the back-end system One or more ActiveMQ long-retry queues One or more ActiveMQ dead-letter queues for manual processing D. One or more On Error scopes to assist calling the back-end system An Until Successful scope containing VM components for long retries A persistent dead-letter VM queue configured in CloudHub	C	

An insurance company is designing a hybrid, load-balanced, single-cluster runtime production environment. Due to performance Service LevelAgreement (SLA) goals, it is looking into running the Mule applications in an active-active multi-node cluster configuration.

What should be considered when running its Mule applications in this type of environment?

A. A Mule application deployed to multiple nodes runs in isolation from the other nodes in the cluster

B. Although the cluster environment is fully installed, configured, and running, it will not process any requests until an outage condition isdetected by the primary node in the cluster

C. All event sources, regardless of type, can be configured as the target source by the primary node in the cluster

D. An external load balancer is required to distribute incoming TCP/HTTP requests throughout the cluster nodes

Time Remaining: 01:27:49 Hide

12 of 60. Why would an Enterprise Architect use a single enterprise-wide canonical data model (CDM) when designing an integration solution using Anypoint Platform?

A. To leverage a data abstraction layer that shields existing Mule applications from non-backward compatible changes to the model's data structure

B. To remove the need to perform data transformation when processing message payloads in Mule applications

C. To automate AI-enabled API implementation generation based on normalized backend databases from separate vendors

D. To reduce dependencies when integrating multiple systems that use different data formats

Mark this item for later review.

48 of 60.

A corporation has deployed Mule applications to different customer-hosted Mule runtimes. Mule applications deployed to these Mule runtimes are managed by Anypoint Platform.

What needs to be installed or configured (if anything) to monitor these Mule applications from Anypoint Monitoring, and how is monitoring data from each Mule application sent to Anypoint Monitoring?

A. ○ Leave the out-of-the-box Anypoint Monitoring agent unchanged in its default Mule runtime installation.
Each Anypoint Monitoring agent sends monitoring data from the Mule applications running in its Mule runtime to Runtime Manager, then Runtime Manager sends monitoring data to Anypoint Monitoring.

B. ◉ Install an Anypoint Monitoring agent on each Mule runtime.
Each Anypoint Monitoring agent sends monitoring data from the Mule applications running in its Mule runtime to Anypoint Monitoring.

C. ○ Install a Runtime Manager agent on each Mule runtime.
Each Runtime Manager agent sends monitoring data from the Mule applications running in its Mule runtime to Runtime Manager, then Runtime Manager sends monitoring data to Anypoint Monitoring.

D. ○ Enable monitoring of individual Mule applications from the Runtime Manager application settings.
Runtime Manager sends monitoring data to Anypoint Monitoring for each deployed Mule application.

23 of 60.

An external API frequently invokes an Employees System API to fetch employee data from a MySQL database. The architect must design a caching strategy to query the database only when there is an update to the Employees table or else return a cached response in order to minimize the number of redundant transactions being handled by the database.

What must the architect do to achieve the caching objective?

A. ○ Use an On Table Row operation configured with the Employees table and call invalidate cache.
Use an object-store-caching-strategy and the default expiration interval.

B. ◉ Use a Scheduler with a fixed frequency set to every hour, triggering an invalidate cache flow.
Use an object-store-caching-strategy and set the expiration interval to 1 hour.

C. ○ Use a Scheduler with a fixed frequency set to every hour to trigger an invalidate cache flow.
Use an object-store-caching-strategy and the default expiration interval.

D. ○ Use an On Table Row operation configured with the Employees table, call invalidate cache, and hardcode the new Employees data to cache.
Use an object-store-caching-strategy and set the expiration interval to 1 hour.

☐ Mark this item for later review.