

## Question 2 of 60

recording



Choose 1 option.

1 org uses a r-nc customer-hosted Mule runtime cluster to host one stateless API implementation. The API is  
ces: TTPS th igh ad balancer that uses round-robin for load distribution. Each node in the cluster has been sized  
acc , mes th irrel umber of requests.

Two nodes in the cluster experience a power outage and are no longer available. The load balancer detects the outage and blocks the two unavailable nodes from receiving further HTTP requests.

What performance-related consequence is guaranteed to happen, on average, assuming the remaining cluster nodes are fully operational?

- ☐ A. 100% increase in the number of requests received by each remaining node
- ☐ B. 100% increase in the average response time of the API implementation
- ☒ C. 50% increase in the number of thread pools on the remaining servers
- ☐ D. 50% reduction in the throughput of the API implementation



Review



Comment

Back

Next

### Question 3 of 60

recording



Choose 1 option.

An external API frequently invokes an Employees System API to fetch employee data from a MySQL database. The architect must design a caching strategy to query the database only when there is an update to the Employees table or else return a cached response in order to minimize the number of redundant transactions being handled by the database.

What must the architect do to achieve the caching objective?

- ☒ A. Use an On Table Row operation configured with the Employees table and call invalidate cache.  
Use an object-store-caching-strategy and the default expiration interval.
- ☐ B. Use a Scheduler with a fixed frequency set to every hour, triggering an invalidate cache flow.  
Use an object-store-caching-strategy and set the expiration interval to 1 hour.
- ☐ C. Use a Scheduler with a fixed frequency set to every hour to trigger an invalidate cache flow.  
Use an object-store-caching-strategy and the default expiration interval.
- ☐ D. Use an On Table Row operation configured with the Employees table, call invalidate cache, and hardcode the new Employees data to cache.  
Use an object-store-caching-strategy and set the expiration interval to 1 hour.



Review



Comment

Back

Next



## Question 4 of 60

Choose 1 option.

A large life sciences customer plans to use the Mule Tracing module with the Mapped Diagnostic Context (MDC) logging operations to enrich logging in its Mule application and to improve tracking by providing more context in the Mule application logs. The customer also wants to improve throughput and lower the message processing latency in its Mule application flows.

After installing the Mule Tracing module in the Mule application, how should logging be performed in flows in Mule applications, and what should be changed in the log4j2.xml files?

- ☐ A. In the flows, add Mule Tracing module Set logging variable operations before any Core Logger components. In log4j2.xml files, change the appender's pattern layout to use %MDC and then assign the appender to a Logger or Root element.
- ☐ B. In the flows, wrap Logger components in Async scopes. In log4j2.xml files, change the appender's pattern layout to use the %asyncLogger placeholder and then assign the appender to a Logger or Root element.
- ☐ C. In the flows, add Mule Tracing module Set logging variable operations before any Core Logger components. In log4j2.xml files, change the appender's pattern layout to use %asyncLogger placeholder and then assign the appender to an AsyncLogger element.
- ☒ D. In the flows, add Mule Tracing module Set logging variable operations before any Core Logger components. In log4j2.xml files, change the appender's pattern layout to use the %MDC placeholder and then assign the appender to an AsyncLogger element.



Review



Comment

Back

Next

## Question 5 of 60

Recording



Choose 1 option.

A retail company is implementing a MuleSoft API to get inventory details from two vendors by invoking each vendor's online applications. Due to network issues, the invocations to the vendor applications are timing out intermittently, but the requests are successful after re-invoking each vendor application.

What is the most performant way of implementing the API to invoke each vendor application and to retry invocations that generate timeout errors?

- ☒ A. Use a Scatter-Gather scope to invoke each vendor application on a separate route.  
Use an Until-Successful scope in each route to retry requests that raise timeout errors.
- ☐ B. Use a Round-Robin scope to invoke each vendor application on a separate route.  
Use a Try-Catch scope in each route to retry requests that raise timeout errors.
- ☐ C. Use a Choice scope to invoke each vendor application on a separate route.  
Place the Choice scope inside an Until-Successful scope to retry requests that raise timeout errors.
- ☐ D. Use a For-Each scope to invoke the two vendor applications in series, one after the other.  
Place the For-Each scope inside an Until-Successful scope to retry requests that raise timeout errors.



Review



Comment

Back

Next

## Question 6 of 60

Recording



Choose 1 option.

An insurance company is implementing a new Mule application that supports a set of critical functions driven by a REST-API-enabled claim payment rules engine hosted on Oracle ERP. The new Mule application requires many data transformation operations as it performs batch processing logic. The company wants to leverage and reuse as many of its existing Java-based capabilities (such as classes, objects, frameworks, and data models) as possible.

What approach should be used when implementing the required data mappings and transformations between other Mule application components and Oracle ERP?

- ☒ A. Transform by calling any suitable Java class from DataWeave
- ☐ B. Create new metadata RAML classes in the Mule runtime corresponding to existing Java objects, and then perform required transformations via DataWeave
- ☐ C. Invoke any of the appropriate Java methods directly, create corresponding metadata RAML classes, and then perform required transformations via DataWeave
- ☐ D. Transform via the XSLT module



Review



Comment

Back

Next



## Question 7 of 60

Choose 1 option.

What requirement prevents using Anypoint MQ as the messaging broker for a Mule application?

- ☒ A. When the messaging broker must be deployed on-premises
- ☐ B. When the payload sent through the message broker must be encrypted
- ☐ C. When the payload sent through the message broker must use XML format
- ☐ D. When the messaging broker must support point-to-point messaging



Review

Comment

Back

Next

## Question 8 of 60



Choose 1 option.

A company is designing a Mule application to consume batch data from a partner's FTPS server, `ftps.partner.com`. The data files have been compressed (but not encrypted) and then digitally signed by the partner using Pretty Good Privacy (PGP).

What inputs are required for the application to securely login to the FTPS server and to consume these files?

- ☐ A. The company's FTPS server login username and password  
A TLS context trust store containing a public certificate for the company  
The company's PGP public key that was used to sign the files
- ☒ B. The company's FTPS server login username and password  
A TLS context trust store containing a public certificate for `ftps.partner.com`  
The partner's PGP public key that was used to sign the files
- ☐ C. The partner's PGP public key used by the company to login to the FTPS server  
A TLS context key store containing the private key for the company  
The partner's PGP private key that was used to sign the files
- ☐ D. The partner's PGP public key used by the company to login to the FTPS server  
A TLS context key store containing the private key for `ftps.partner.com`  
The company's PGP private key that was used to sign the files



Review



Comment

Back

Next



## Question 9 of 60

Choose 1 option.

An application deployed in a CloudHub production environment is throwing intermittent exceptions. To troubleshoot the issue, a developer wants to enable DEBUG level logging in production for the `org.apache.cxf` package in the deployed Mule application.

How should the developer update the logging configuration in order to enable this package specific debugging?

- ☐ A. In the Mule application's `log4j2.xml` file, change the root logger's level property to `DEBUG`, then redeploy the Mule application to the CloudHub production environment
- ☐ B. In Anypoint Runtime Manager, in the Deployed Application Properties tab for the Mule application, add a line item with `DEBUG` level for package `org.apache.cxf` and apply the changes
- ☒ C. In the Mule application's `log4j2.xml` file, add an `AsyncLogger` element with name property set to `org.apache.cxf` and level set to `DEBUG`, then redeploy the Mule application in the CloudHub production environment
- ☐ D. In Anypoint Monitoring, define a logging search query with class property set to `org.apache.cxf` and level set to `DEBUG`



Review



Comment

Back

Next

## Question 10 of 60

recording



Choose 1 option.

A manufacturing company wants to share inventory updates with dealers D1 and D2 asynchronously and concurrently via Queues Q1 and Q2. Dealer D1 must consume the messages from the Queue Q1, and Dealer D2 must consume the messages from the Queue Q2.

Dealer D1 has implemented a retry mechanism to reprocess the transaction in case of any errors while processing the inventory updates. Dealer D2 has not implemented any retry mechanism.

How should the dealers acknowledge the message to avoid message loss and minimize impact on the current implementation?

- ☒ A. Dealer D1 can use AUTO acknowledgement and Dealer D2 must use MANUAL acknowledgment and acknowledge the message after successful processing
- ☐ B. Dealer D1 can use AUTO acknowledgement and Dealer D2 can use IMMEDIATE acknowledgment and acknowledge the message after successful processing
- ☐ C. Dealer D1 and Dealer D2 must use AUTO acknowledgement and acknowledge the message after successful processing
- ☐ D. Dealer D2 must use AUTO acknowledgement and Dealer D1 can use MANUAL acknowledgment and acknowledge the message after successful processing



Review



Comment

Back

Next

## Question 11 of 60



Choose 1 option.

A corporation has deployed Mule applications to different customer-hosted Mule runtimes. Mule applications deployed to these Mule runtimes are managed by Anypoint Platform.

What needs to be installed or configured (if anything) to monitor these Mule applications from Anypoint Monitoring, and how is monitoring data from each Mule application sent to Anypoint Monitoring?

- ☐ A. Enable monitoring of individual Mule applications from the Runtime Manager application settings.  
Runtime Manager sends monitoring data to Anypoint Monitoring for each deployed Mule application.
- ☒ B. Install an Anypoint Monitoring agent on each Mule runtime.  
Each Anypoint Monitoring agent sends monitoring data from the Mule applications running in its Mule runtime to Anypoint Monitoring.
- ☐ C. Leave the out-of-the-box Anypoint Monitoring agent unchanged in its default Mule runtime installation.  
Each Anypoint Monitoring agent sends monitoring data from the Mule applications running in its Mule runtime to Runtime Manager, then Runtime Manager sends monitoring data to Anypoint Monitoring.
- ☐ D. Install a Runtime Manager agent on each Mule runtime.  
Each Runtime Manager agent sends monitoring data from the Mule applications running in its Mule runtime to Runtime Manager, then Runtime Manager sends monitoring data to Anypoint Monitoring.



Review



Comment

Back

Next

## Question 12 of 60

Recording



Choose 1 option.

When a Mule application using VM queues is deployed to a customer-hosted cluster or multiple CloudHub v1.0 workers/replicas, how are messages consumed across the nodes?

- ☐ A. Sequentially, only from the primary node
- ☒ B. In a non-deterministic way
- ☐ C. Round-robin, within an XA transaction
- ☐ D. Sequentially, from a dedicated Anypoint MQ queue



Review



Comment

Back

Next

## Question 13 of 60



Choose 1 option.

An organization plans to use the Anypoint Platform audit logging service to log Anypoint MQ actions.

What consideration must be kept in mind when leveraging Anypoint MQ Audit Logs?

- ☐ A. Anypoint MQ Audit Logs include logs for failed Anypoint MQ operations
- ☐ B. Anypoint MQ Audit Logs include logs for queue create, delete, modify, and purge operations
- ☒ C. Anypoint MQ Audit Logs include logs for sending, receiving, or browsing messages



Review



Comment

Back

Next

## Question 14 of 60



Choose 1 option.

A healthcare provider is implementing Anypoint Platform to manage its application infrastructure and is using the customer-hosted runtime for its business due to certain financial requirements it must meet. It has built a number of synchronous APIs and is currently hosting these on a Mule runtime on one server. These applications make use of a number of components, including heavy use of Object Stores and VM Queues.

Business has grown rapidly in the last year, and the healthcare provider is starting to receive reports of reliability issues from its applications. The DevOps team indicates that the APIs are currently handling too many requests, and this is overloading the server. The team has also mentioned that there is significant downtime when the server is down for maintenance.

What can the team add or upgrade to mitigate these issues?

- ☐ A. Add additional servers in a server group configuration
- ☐ B. Upgrade the networking hardware of the server
- ☐ C. Upgrade the CPU and memory of the server
- ☒ D. Add a load balancer and add additional servers in a cluster configuration



Review



Comment

Back

Next

## Question 15 of 60



Choose 1 option.

A bank is implementing a REST API in a Mule application to receive an array of accounts from an online banking platform user interface (UI), retrieve account balances for those accounts from a backend Finance system, and then return the account balances so they can be displayed in the online banking platform UI. As part of the processing, the MuleSoft API also needs to insert the retrieved account data into an Audit Database for auditing purposes. the auditing process should not add latency to the account balance retrieval response back to the online banking platform UI.

The retrieveBalances flow in the Mule application is designed to use an operation in a connector to the Finance system (the Finance operation) that can only look up one account record at a time, and a operation from a different connector to the Audit system (the Audit operation) that can only insert one account record at a time.

To best meet the performance-related requirements, what scope or scopes should be used and how should they be used to incorporate the Finance operation and Audit operation into the retrieveBalances flow?

- ☐ A. Wrap the Finance operation in a Until-Successful scope.  
Wrap the Audit operation in a Try-Catch scope.
- ☒ B. Wrap both connector operations in a For-Each scope.
- ☐ C. Wrap the Finance operation in a Parallel For-Each scope.  
Wrap the Audit operation in a Async scope.
- ☐ D. Wrap both connector operations in a Async scope.



Review



Comment

Back

Next

## Question 16 of 60

Recording



Choose 1 option.

An organization's security requirements mandate centralized control at all times over authentication and authorization of external applications when invoking web APIs managed on Anypoint Platform.

What Anypoint Platform feature is the most straightforward and maintainable to use to meet this requirement?

- ☐ A. Identity management configured in Access Management
- ☐ B. Enterprise Security module coded in Mule applications
- ☒ C. Client management configured in Access Management
- ☐ D. External access configured in API Manager



Review



Comment

Back

Next

## Question 17 of 60



Choose 1 option.

A Mule application contains a Batch Job scope with several Batch Step scopes. The Batch Job scope is configured with a batch block size of 25. A payload with 4,000 records is received by the Batch Job scope.

When there are no errors, how does the Batch Job scope process records within and between the Batch Step scopes?

- ☐ A. The Batch Job scope processes multiple record blocks in parallel, and a block of 25 records can jump ahead to the next Batch Step scope over an earlier block of records.  
Each Batch Step scope is invoked with one record in the payload of the received Mule event.  
For each Batch Step scope, all 25 records within a block are processed in parallel.  
All the records in a block must be completed before the block of 25 records is available to the next Batch Step scope.
- ☒ B. The Batch Job scope processes multiple record blocks in parallel, and a block of 25 records can jump ahead to the next Batch Step scope over an earlier block of records.  
Each Batch Step scope is invoked with one record in the payload of the received Mule event.  
For each Batch Step scope, all 25 records within a block are processed sequentially.  
All the records in a block must be completed before the block of 25 records is available to the next Batch Step scope.
- ☐ C. The Batch Job scope processes multiple record blocks in parallel.  
Each Batch Step scope is invoked with a batch of 25 records in the payload of the received Mule event.  
For each Batch Step scope, all 4,000 records are processed in parallel.  
Individual records can jump ahead to the next Batch Step scope before the rest of the records finish processing



Review



Comment

Back

Next

## Question 18 of 60

Recording



Choose 1 option.

An organization plans to extend its Mule APIs to the EU (Frankfurt) region. Currently, all Mule applications are deployed to CloudHub 1.0 in the default North American region, from the North America control plane, following this naming convention: **{API-name}-{environment}** (for example, Orders-sapi-dev, Orders-sapi-qa, Orders-sapi-prod, etc.).

There is no network restriction to block communications between APIs.

What strategy should be implemented in order to deploy the same Mule APIs to the CloudHub 1.0 EU region from the North America control plane, as well as to minimize latency between APIs and target users and systems in Europe?

- ☐ A. In API Manager, leave the Region property blank (default) to deploy an API proxy named **{API-name}-proxy-{environment}.de-c1** for each Mule application.

Communicate the new url **{API-name}-proxy-{environment}.de-c1.cloudhub.io** to the consuming API clients in Europe.

- ☐ B. In API Manager, set the Region property to EU (Frankfurt) to create an API proxy named **{API-name}-proxy-{environment}** for each Mule application.

Communicate the new url **{API-name}-proxy-{environment}.de-c1.cloudhub.io** to the consuming API clients in Europe.

- ☒ C. In Runtime Manager, for each Mule application deployment, set the Region property to EU (Frankfurt) and reuse the same Mule application name as in the North American region.



Review



Comment

Back

Next

## Question 18 of 60

Recording



Choose 1 option.

- ☐ A. In API Manager, leave the Region property blank (default) to deploy an API proxy named **{API-name}-proxy-{environment}.de-c1** for each Mule application.

Communicate the new url **{API-name}-proxy-{environment}.de-c1.cloudhub.io** to the consuming API clients in Europe.

- ☒ B. In API Manager, set the Region property to EU (Frankfurt) to create an API proxy named **{API-name}-proxy-{environment}** for each Mule application.

Communicate the new url **{API-name}-proxy-{environment}.de-c1.cloudhub.io** to the consuming API clients in Europe.

- ☐ C. In Runtime Manager, for each Mule application deployment, set the Region property to EU (Frankfurt) and reuse the same Mule application name as in the North American region.

Communicate the new urls **{API-name}-{environment}.de-c1.cloudhub.io** to the consuming API clients in Europe.

- ☐ D. In Runtime Manager, for each Mule application deployment, leave the Region property blank (default) and change the Mule application name to **{API-name}-{environment}.de-c1**.

Communicate the new urls **{API-name}-{environment}.de-c1.cloudhub.io** to the consuming API clients in Europe.



Review



Comment

Back

Next

## Question 19 of 60

recording



Choose 1 option.

A Mule application is designed to call the Google Maps API to perform a distance computation, and is deployed to CloudHub.

At a minimum, what (if anything) should be configured in the TLS context of the HTTP Request configuration to meet these requirements?

- ☒ A. The configuration is built-in and nothing extra is required for the TLS context
- ☐ B. Download the Google public certificate from a browser, generate a JKS file from it, and add it in TrustStore as part of the TLS context
- ☐ C. Download the Google public certificate from a browser, generate a JKS file from it, and add it in KeyStore as part of the TLS context
- ☐ D. Request a private key from Google, create a PKCS#12 file with it, and add it in KeyStore as part of the TLS context



Review



Comment

Back

Next

## Question 20 of 60

Recording



Choose 1 option.

A leading eCommerce giant will use MuleSoft APIs on Runtime Fabric (RTF) to process customer orders. Some customer-sensitive information, such as credit card information, is required in request payloads or is included in response payloads in some of the APIs. Other API requests and responses are not authorized to access some of this customer-sensitive information but have been implemented to validate and transform based on the structure and format of this customer-sensitive information (such as account IDs, phone numbers, and postal codes).

What approach configures an API gateway to hide sensitive data exchanged between API consumers and API implementations, but can convert tokenized fields back to their original value for other API requests or responses, without having to recode the API implementations?

- ☐ A. Create a masking format and use it to apply a tokenization policy in an API gateway to mask sensitive values in message payloads with \*\*\*\* characters, and apply a corresponding detokenization policy to return the original values to other APIs
- ☐ B. Use a field-level encryption policy in an API gateway to replace sensitive fields in message payload with encrypted values, and apply a corresponding field-level decryption policy to return the original values to other APIs
- ☐ C. Create both masking and tokenization formats and use both to apply a tokenization policy in an API gateway to mask sensitive values in message payloads with \*\*\*\* characters, and apply a corresponding detokenization policy to return the original values to other APIs
- ☒ D. Create a tokenization format and use it to apply a tokenization policy in an API gateway to replace sensitive



Review



Comment

Back

Next

## Question 21 of 60

Recording



Choose 1 option.

An architect is designing a Mule application to meet the following two requirements:

1. The application must process files asynchronously and reliably from an FTPS server to a back-end database using VM intermediary queues for load-balancing Mule events.
2. The application must process a medium rate of records from a source to a target system using a Batch Job scope.

To make the Mule application more reliable, the Mule application will be deployed to two CloudHub 1.0 workers.

Following MuleSoft-recommended best practices, how should the Mule application deployment typically be configured in Runtime Manager to best support the performance and reliability goals of both the Batch Job scope and the file processing VM queues?

- ☒ A. Check the Persistent VM queues checkbox in the application deployment configuration
- ☐ B. In the Runtime Manager Properties tab, disable persistent VM queues for Batch Job scopes
- ☐ C. In the Runtime Manager Properties tab, enable persistent VM queues for the FTPS connector
- ☐ D. Check the Non-persistent VM queues checkbox in the application deployment configuration



Review



Comment

Back

Next

## Question 22 of 60

Recording



Choose 1 option.

An organization has strict unit test requirements that mandate every Mule application must have an MUnit test suite with a test case defined for each flow and a minimum test coverage of 80 percent.

A developer is building an MUnit test suite for a newly developed Mule application that sends API requests to an external REST API.

What is an effective approach for successfully executing the MUnit tests of this Mule application while still achieving the required test coverage for the MUnit tests?

- ☐ A. Mock the REST API invocations in the MUnit tests, and then forward mocked requests to the external REST API
- ☒ B. Mock the REST API invocations in the MUnit tests, and then return a mock response for those invocations
- ☐ C. Create a mocking service flow to simulate standard responses from the REST API, and then configure the Mule flows to call the mocking service flow
- ☐ D. Invoke the external endpoint of the REST API from the Mule flows



Review



Comment

Back

Next

## Question 23 of 60

Recording



Choose 1 option.

A project team is working on an API implementation using an API RAML specification as a starting point. The team has updated the RAML specification to include new operations, and has published a new version to Exchange. Meanwhile, another team is working on an API client Mule application to consume the same API implementation.

During this development, what must be performed by the Mule application team to take advantage of the newly added operations?

- ☐ A. Scaffold the API client application with APIkit using the new RAML specification
- ☒ B. Update the REST connector from Exchange in the API client application
- ☐ C. Update the API connector in the API implementation and publish to Exchange
- ☐ D. Copy the API implementation's updated APIkit component into the API client



Review



Comment

Back

Next

## Question 24 of 60




Choose 1 option.

Refer to the exhibit.

One of the back-end systems invoked by the API implementation enforces rate limits on the number of requests a particular client can make. Both the back-end system and the API implementation are deployed to several non-production environments, including the staging environment, and to a production environment.

Rate limiting of the back-end system applies to all non-production environments. The production environment, however, does not have any rate limiting.

What is the most effective approach to conduct performance tests of the API implementation in the non-production staging environment?




Review



Comment

Back

Next



- ☐ A. Conduct scaled-down performance tests in the staging environment against the rate-limited back-end system. Then, upscale performance results to full production scale.
- ☐ B. Include logic within the API implementation that bypasses invocations of the back-end system in the staging environment, and invoke a mocking service that replicates typical back-end system responses. Then, conduct performance tests using this API implementation.
- ☒ C. Create a mocking service that replicates the back-end system's production performance characteristics. Then, configure the API implementation to use the mocking service and conduct the performance tests.
- ☐ D. Use MUnit to simulate standard responses from the back-end system. Then, conduct performance tests to identify other bottlenecks in the system.



Review



Comment

Back

Next

Choose 1 option.

A cluster with two customer-hosted Mule runtimes is hosting an application that has a flow with a JMS listener configured to subscribe to messages from a topic destination.

How should the JMS listener's `primaryNodeOnly` parameter and subscription parameter be configured so the JMS listener can receive messages in all the nodes of the cluster but only process each message once?

- ☒ A. Use the parameter `primaryNodeOnly="false"` on the JMS listener with a shared subscription
- ☐ B. Use the parameter `primaryNodeOnly="true"` on the JMS listener with a non-shared subscription
- ☐ C. Use the parameter `primaryNodeOnly="true"` on the JMS listener with a shared subscription
- ☐ D. Use the parameter `primaryNodeOnly="false"` on the JMS listener with a non-shared subscription



Review



Comment

Back

Next

Choose 1 option.

A gaming company has implemented an API as a Mule application and deployed the API implementation to a CloudHub 2.0 private space. The API implementation must connect to a mainframe application running in the customer's on-premises corporate data center and also to a Kafka cluster running in an Amazon AWS VPC.

What is the most efficient way to enable the API to securely connect from its private space to the mainframe application and Kafka cluster?

- ☒ A. In Runtime Manager, configure an Anypoint VPN to route from the CloudHub 2.0 private space to the on-premises data center. In the AWS account, attach the private space to an AWS transit gateway that routes to the AWS VPC.



- ☐ B. In Runtime Manager, set up VPC peering between the CloudHub 2.0 private network and the on-premises data center. In the AWS account, set up VPC peering between the AWS VPC and the CloudHub 2.0 private network.



Review



Comment

Back

Next


Choose 1 option.

VPC Peering      AWS VPC

- ☐ C. In the AWS account, attach the private space directly to the AWS VPC.  
In the AWS account, use an AWS transit gateway to route from the AWS VPC to the on-premises data center.



- ☐ D. In the AWS account, attach the CloudHub 2.0 private space to an AWS transit gateway that routes from the CloudHub 2.0 private space to the on-premises data center. In Runtime Manager, configure an Anypoint VPN to route from the CloudHub 2.0 private space to the AWS VPC.



Review



Comment

Back

Next

Choose 1 option.

A manufacturing company plans to deploy Mule applications to its own Azure Kubernetes service infrastructure. The organization wants to make the Mule applications more available and robust by deploying each Mule application to an isolated Mule runtime in a Docker container while managing all the Mule applications from the MuleSoft-hosted control plane.

What choice of runtime plane meets these organizational requirements?

- ☐ A. CloudHub 2.0
- ☒ B. Anypoint Runtime Fabric
- ☐ C. Anypoint Service Mesh
- ☐ D. Customer-hosted self-provisioned runtime plane



Review



Comment

Back

Next



Choose 1 option.

An organization is successfully using API-led connectivity. However, as the application network grows, all the manually performed tasks to publish, share and discover, register, apply policies to, and deploy an API are becoming repetitive, which is driving the organization to automate this process using efficient CI/CD pipelines.

Considering Anypoint Platform's capabilities, how should the organization automate its API lifecycle?

- ☐ A. Use Anypoint CLI or Anypoint Platform REST APIs with a scripting language such as Groovy
- ☐ B. Use Runtime Manager REST APIs for API management and Maven for API deployment
- ☐ C. Use Maven with a custom configuration required for the API lifecycle
- ☒ D. Use Exchange REST APIs for API management and Maven for API deployment



Review



Comment

Back

Next

Choose 1 option.

An organization is using MuleSoft to develop APIs and automatically deploy them to both CloudHub and on-premises targets. Recently, it has decided to include Runtime Fabric deployment targets as well, and Mule infrastructure is set up for this option.

What can now be used to automate deploying Mule applications to all three types of deployment targets?

- ☐ A. Anypoint Monitoring agent configuration file
- ☐ B. The Object Store v2 REST API
- ☐ C. The Runtime Manager agent configuration file
- ☒ D. The Mule Maven plugin



Review



Comment

Back

Next

## Question 30 of 60

Recording



Choose 1 option.


Refer to the exhibits.

An Order microservice and a Fulfillment microservice are being designed to communicate with their clients through message-based integration (and NOT through API invocations).

The Order microservice publishes an Order message (a kind of command message) containing the details of an order to be filled. The intent is that Order messages are only consumed by one Mule application: the Fulfillment microservice.

The Fulfillment microservice consumes Order messages, fills the orders described therein, and then publishes a Fulfillment message (a kind of event message). Each Fulfillment message can be consumed by any interested Mule application, and the Order microservice is one such Mule application.

What is the recommended choice of message broker(s) and message destination(s) in this scenario?



- ☐ A. Order messages are sent to an Anypoint MQ exchange. Fulfillment messages are sent to an Anypoint MQ queue. Both microservices interact with Anypoint MQ as the message broker, which must therefore scale to support the



Review



Comment

Back


Next

Question 30 of 60


Recording

Choose 1 option.

- ☒ B. Order messages are sent to a JMS topic.  
Fulfillment messages are sent to a JMS queue.  
The Order microservice interacts with one JMS provider (message broker) and the Fulfillment microservice interacts with a different JMS provider so that both message brokers can be chosen and scaled to best support the load of each microservice.



- ☐ C. Order messages are sent to a JMS queue.  
Fulfillment messages are sent to a JMS topic.  
Both microservices interact with the same JMS provider (message broker) instance, which must therefore scale to support the load of both microservices.



Review



Comment

Back

Next

Choose 1 option.

An organization is in the process of building automated deployments using a CI/CD process. As a part of automated deployments, it wants to apply policies to API instances.

What tool can the organization use to promote and deploy API Manager policies?

- ☐ A. Runtime Manager agent
- ☐ B. Mule Maven plugin
- ☒ C. Anypoint CLI
- ☐ D. MUnit Maven plugin



Review



Comment

Back

Next

A healthcare client must connect its Mule APIs that are deployed to CloudHub 1.0 to systems in an existing on-premises datacenter through private and secure connections.

Which network architecture supports these security requirements to connect the Mule APIs to the on-premises datacenter?

- ☐ A. Static IP addresses for the Mule applications deployed to the CloudHub 1.0 Shared Worker Cloud  
Matching IP include listing in the load balancer and firewall rules in the VPC and on-premises network
- ☐ B. Static IP addresses for the Mule applications deployed to the CloudHub 1.0 Shared Worker Cloud  
Matching firewall rules and IP include listing in the on-premises network
- ☐ C. An Anypoint VPC connected to the on-premises network using a dedicated load balancer  
Matching IP include listing in the load balancer and firewall rules in the VPC and on-premises network
- ☒ D. An Anypoint VPC connected to the on-premises network using an IPsec tunnel  
Matching firewall rules in the VPC and on-premises network



Review



Comment

Back

Next

Choose 1 option.

Why would an Enterprise Architect use a single enterprise-wide canonical data model (CDM) when designing an integration solution using Anypoint Platform?

- ☐ A. To remove the need to perform data transformation when processing message payloads in Mule applications
- ☐ B. To automate AI-enabled API implementation generation based on normalized backend databases from separate vendors
- ☐ C. To leverage a data abstraction layer that shields existing Mule applications from non-backward compatible changes to the model's data structure
- ☒ D. To reduce dependencies when integrating multiple systems that use different data formats



Review



Comment

Back

Next

Choose 1 option.

A payment processing company has implemented a Payment Processing API Mule application to process credit card and debit card transactions. Because the Payment Processing API handles highly sensitive information, the payment processing company requires that data must be encrypted both in-transit and at-rest.

To meet these security requirements, consumers of the Payment Processing API must create request message payloads in a JSON format specified by the API, and the message payload values must be encrypted.

How can the Payment Processing API validate requests received from API consumers?

- ☒ A. A Transport Layer Security (TLS) - Inbound policy can be applied in API Manager to decrypt the message payload and the Mule application implementation can then use the JSON Validation module to validate the JSON data
- ☐ B. The Mule application implementation can use the APIkit module to decrypt and then validate the JSON data
- ☐ C. The Mule application implementation can use the Validation module to decrypt and then validate the JSON data
- ☐ D. The Mule application implementation can use DataWeave to decrypt the message payload and then use the JSON Scheme Validation module to validate the JSON data



Review



Comment

Back

Next

Choose 1 option.

An organization has deployed Runtime Fabric on a single server configured with an eight-node cluster using the performance profile. An API uses a non-persistent Object Store for maintaining state data from deployed Mule applications.

What will be the impact to the state data if the server crashes?

- ☐ A. State data is preserved
- ☒ B. State data is lost
- ☐ C. State data is rolled back to a previously saved version
- ☐ D. State data is duplicated



Review



Comment

Back

Next

Choose 1 option.

An organization is building out a test suite for their application using MUnit. The Integration Architect has recommended using Test Recorder in Anypoint Studio to record the processing flows and then configure unit tests based on the captured events.

What is a core consideration that must be kept in mind while using Test Recorder?

- ☒ A. Tests for flows cannot be created if Mule errors are raised inside the flows, even if the errors are handled by On-Error Continue error handlers
- ☐ B. The Recorder supports mocking a message before or inside a Foreach processor
- ☐ C. The Recorder supports loops where the structure of the data being tested changes inside the iteration
- ☐ D. Mocking values resulting from parallel processes are possible and will not affect the execution of the processors that follow in the test



Review



Comment

Back

Next

Choose 1 option.

An organization has decided on a cloud migration strategy to minimize the organization's own IT resources. Currently, the organization has all of its Mule applications running on its own premises and uses an on-premises load balancer that exposes all APIs under the base URL `https://api.acme.com`.

As part of the migration strategy, the organization plans to migrate all of its Mule applications and the on-premises load balancer to CloudHub 1.0.

What is the most straightforward and cost-effective approach to Mule application deployment and load balancing that preserves the public URLs?

- ☐ A. Deploy the Mule applications to CloudHub 1.0.  
Update the CNAME record for `api.acme.com` in the organization's DNS server to point to the A record of the CloudHub Shared Load Balancer (SLB).  
Apply mapping rules in the SLB to map URLs to their corresponding Mule applications.
- ☒ B. Deploy the Mule applications to CloudHub 1.0.  
Update the CNAME record for `api.acme.com` in the organization's DNS server to point to the A record of a CloudHub 1.0 Dedicated Load Balancer (DLB).  
Apply mapping rules in the DLB to map URLs to their corresponding Mule applications.
- ☐ C. For each migrated Mule application, deploy an API proxy Mule application to CloudHub 1.0, with all traffic to the Mule applications routed through a CloudHub 1.0 Dedicated Load Balancer (DLB).  
Update the CNAME record for `api.acme.com` in the organization's DNS server to point to the A record of the CloudHub 1.0 DLB.  
Apply mapping rules in the DLB to map each API proxy application to its corresponding Mule application.
- ☐ D. Deploy the Mule applications to CloudHub 1.0.  
Update the CNAME record for `api.acme.com` in the organization's DNS server to point to the A record of a CloudHub 1.0 Shared Load Balancer (SLB).  
Apply mapping rules in the SLB to map URLs to their corresponding Mule applications.



Review



Comment

Back

Next

Choose 1 option.

An organization has built a large monolithic application over the years and is currently looking to transition to a microservices architecture. During this transition phase, it has built some System APIs as Mule applications that provide access to different pieces of functionality of the monolithic application. The System APIs are deployed to Mule runtimes in a customer-hosted EC2 instance in AWS.

The System APIs communicate directly with one database used by the monolithic application. Currently, a Database connector and configuration for that connector must be built into every project. There is a plan to partition this database into various smaller data stores in the future.

What is the MuleSoft-recommended best practice to share the connector and configuration information among the APIs?

- ☐ A. Build another System API that connects to the database, and refactor all the other APIs to make requests through the new System API to access the database
- ☐ B. Create an API proxy for each System API and share the Database connector configuration with all the API proxies via an automated policy
- ☐ C. Build a separate Mule domain project for each API, and configure each of them to use a file on a shared file store to load the configuration information dynamically
- ☒ D. Build a Mule domain project, add the Database connector and configuration to it, and reference this one domain project from each System API



Review



Comment

Back

Next

Choose 1 option.

Which Salesforce API is invoked to deploy, retrieve, create, update, or delete customization information, such as custom object definitions using Mule Salesforce Connectors in a Mule application?

- ☐ A. sObject Platform Action API
- ☐ B. Process Rules API
- ☐ C. User Interface API
- ☒ D. Metadata API



Review 

Comment

Back

Next

Choose 1 option.

Which operation can be performed through a JMX agent enabled in a Mule application?

- ☐ A. Set a particular log4j2 log level to TRACE
- ☒ B. View object store entries
- ☐ C. Deploy a Mule application
- ☐ D. Replay an unsuccessful message



Review



Comment

Back

Next

### Question 41 of 60

Recording



Choose 2 options.


Refer to the exhibit.

A manufacturing company is developing a new set of APIs for its retail business. One of the APIs is a Master Look Up API, which is a System API. The API uses a persistent object-store. This API will be used by almost all other APIs to provide master lookup data.

The Master Look Up API is deployed on two CloudHub workers of 0.1 vCore each because there is a lot of master data to be cached. Most of the master lookup data is stored as a key-value pair. The cache gets refreshed if the key is not found in the cache.

During performance testing, it was determined that the Master Look Up API has a high response time due to the latency of database queries executed to fetch the master lookup data.

What two methods can be used to resolve these performance issues?



- ☐ A. Implement locking to synchronize access to the Object Store
- ☒ B. Upgrade the vCore size from 0.1 vCore to 0.2 vCore
- ☒ C. Implement the HTTP caching policy for all GET endpoints for the Master Look Up API



Review



Comment

Back

Next

Choose 1 option.

A new Mule application has been deployed through Runtime Manager to CloudHub 1.0 using a CI/CD pipeline with sensitive properties set as cleartext. The Runtime Manager Administrator opened a high priority incident ticket about this violation of their security requirements indicating these sensitive properties values must not be stored or visible in Runtime Manager but should be changeable in Runtime Manager by Administrators with proper permissions.

How can the Mule application be deployed while safely hiding the sensitive properties?

- ☐ A. Create a variable for each sensitive property and declare them as hidden in the CI/CD pipeline scripts
- ☐ B. Add a new `wrapper.java.additional.XX` parameter for each sensitive property in the `wrapper.conf` file used by the CI/CD pipeline scripts
- ☒ C. Add encrypted versions of the sensitive properties as global configuration properties in the Mule application
- ☐ D. Add an `ArrayList` of all the sensitive properties' names in the `mule-artifact.json` file of the application



Review



Comment

Back

Next

Question 43 of 60

Recording



Choose 1 option.

A Mule application uses an HTTP Request operation to invoke an external API. The external API follows the HTTP specification for proper status code usage.

What is a possible cause when a 3xx status code is returned to the HTTP Request operation from the external API?

- ☐ A. The request was received and processed by the external API
- ☐ B. The request was not received by the external API
- ☒ C. The request was redirected to a different URL by the external API
- ☐ D. The request was received and rejected by the external API



Review



Comment

Back

Next

Question 44 of 60

Recording



Choose 1 option.

A Mule application is required to periodically process a large data set from a back-end database to Salesforce CRM using a Batch Job scope configured properly to process the high rate of records. The Mule application is deployed to two CloudHub 1.0 workers with the persistent VM queues option disabled (unchecked).

What is the consequence if a worker that is processing records from a Batch Job scope crashes during records processing?

- ☐ A. About half the remaining records will be processed by the second worker and the rest will be lost, resulting in some unprocessed records
- ☒ B. Remaining records will be lost, resulting in unprocessed records
- ☐ C. All the remaining records will be processed by a new replacement worker, resulting in "exactly once" processing of every record
- ☐ D. The entire batch job will be processed from scratch by the second worker, resulting in some duplicate processing



Review



Comment

Back

Next

Question 45 of 60

Recording



Choose 1 option.

A marketing organization is designing a Mule application to process campaign data. The Mule application will periodically check for a file in an SFTP location and process the records in the file. The size of the file can vary from 10MB to 5GB. Due to the limited availability of vCores, the Mule application is deployed to a single CloudHub worker/replica configured with vCore size 0.2.

The application must transform and send different formats of this file to three different downstream SFTP locations.

What is the most direct way to configure the SFTP operations or event sources to process the large files to support these deployment requirements?

- ☐ A. Use a file-stored repeatable stream
- ☐ B. Use an in-memory nonrepeatable stream
- ☒ C. Use a file-stored nonrepeatable stream
- ☐ D. Use an in-memory repeatable stream



Review



Comment

Back

Next

Choose 1 option.

An airline is architecting an API-led connectivity project to integrate its flight data into an online aggregation website. The connections must allow for secured communication, high performance, and asynchronous message exchange.

What are suitable interface technologies for the connectors in this integration assuming that MuleSoft fully supports those technologies and that Anypoint Connectors exist for these interfaces?

- ☒ A. AsyncAPI over SQS  
AMQP with RabbitMQ  
JSON/REST over HTTPS
- ☐ B. XML over ActiveMQ  
XML over FTP  
XML/REST over HTTPS
- ☐ C. SOAP over HTTPS  
IIOP over TLS  
gRPC over HTTPS
- ☐ D. CSV over FTP  
YAML over TLS  
JSON over HTTPS



Review



Comment

Back

Next

Choose 1 option.

In a Mule 4 application that will be deployed to customer-hosted Mule runtimes, the Java module's Invoke Static operation needs to invoke a static method in a dependency JAR file.

How can the dependency be made visible to the Java module's classloader?

- ☐ A. Use a Maven command to include the dependency JAR file when packaging the application
- ☐ B. In the Java Invoke Static module configuration, configure the path and name of the dependency JAR file
- ☐ C. Add the dependency JAR file to the Java classpath by setting JVM parameters
- ☒ D. Configure the dependency as a shared library in the project's POM file



Review



Comment

Back

Next

Choose 1 option.

Following MuleSoft's recommended best practices for API governance and API policies, a project team has used RAML specifications to document and publish functional requirements and detailed design definitions of its APIs. These API specifications have been used by various stakeholders to implement APIs.

Later, the project team requires all API specifications to be augmented with an additional non-functional requirement (NFR) to protect the backend services from a high rate of requests, according to defined service-level agreements (SLAs). The NFR's SLAs are based on a new tiered subscription level "Gold", "Silver", or "Platinum" that must be tied to a new parameter that is being added to the Accounts object in their enterprise data model.

Following MuleSoft's recommended best practices, how should the project team now convey the necessary non-functional requirement to stakeholders?

- ☒ A. Update each API specification with a shared RAML fragment required to implement the NFR and publish the RAML fragment and each modified API specification to Exchange
- ☐ B. Create a shared RAML fragment required to implement the NFR, list each API implementation endpoint in the RAML fragment, and publish the RAML fragment to Exchange
- ☐ C. Create and deploy API proxies in API Manager for the NFR, change the baseUrl in each API specification to the corresponding API proxy implementation endpoint, and publish each modified API specification to Exchange
- ☐ D. Update each API specification with comments about the NFR's SLAs and publish each modified API specification to Exchange



Review



Comment

Back

Next

Choose 1 option.

A Mule application is synchronizing customer data between two different database systems.

What is a main benefit of using eXtended Architecture (XA) transactions rather than local transactions to synchronize these two different database systems?

- ☒ A. An XA transaction automatically rolls back operations against both database systems if any operation fails
- ☐ B. An XA transaction synchronizes the database systems with the least amount of Mule configuration or coding
- ☐ C. An XA transaction handles the largest number of requests in the shortest time
- ☐ D. An XA transaction writes to both database systems quickly and simultaneously



Review



Comment

Back

Next

Choose 1 option.

An organization is designing multiple Mule applications to run on CloudHub in a single Anypoint VPC and that must share data using a common persistent Anypoint Object Store v2 (OSv2).

What design gives these Mule applications access to the same Object Store instance?

- ☐ A. The Object Store v2 can be shared across CloudHub applications with a configured OS v2 connector
- ☒ B. The Object Store v2 REST API configured to access the persistent Object Store
- ☐ C. A VM connector configured to directly access the persistence queue of the persistent Object Store
- ☐ D. An Anypoint MQ connector configured to directly access the persistent Object Store



Review



Comment

Back

Next

Choose 1 option.

A company is designing an integration Mule application to process orders by submitting them to a back-end system for offline processing. Each order will be received by the Mule application through an HTTPS POST and must be acknowledged immediately. Once acknowledged, the order will be submitted to a back-end system. Orders that cannot be successfully submitted due to rejections from the back-end system will need to be processed manually (outside the back-end system).

The Mule application will be deployed to a customer-hosted runtime and will be able to use an existing ActiveMQ broker if needed. The ActiveMQ broker is located inside the organization's firewall. The back-end system has a track record of unreliability due to both minor network connectivity issues and longer outages.

What combination of Mule application components and ActiveMQ queues are required to ensure automatic submission of orders to the back-end system while supporting but minimizing manual order processing?

- ☒ A. An Until Successful scope to call the back-end system  
One or more ActiveMQ long-retry queues  
One or more ActiveMQ dead-letter queues for manual processing
- ☐ B. One or more On Error scopes to assist calling the back-end system  
An Until Successful scope containing VM components for long retries  
A persistent dead-letter VM queue configured in CloudHub
- ☐ C. One or more On Error scopes to assist calling the back-end system  
One or more ActiveMQ long-retry queues  
A persistent dead-letter Object Store configured in the CloudHub Object Store service
- ☐ D. A Batch Job scope to call the back-end system  
An Until Successful scope containing Object Store components for long retries  
A dead-letter object store configured in the Mule application



Review



Comment

Back

Next

Choose 1 option.

A Mule solution implements an Experience API and System API as Mule applications that process messages in this order:

1. The Experience API receives a request Anypoint MQ message (REQU) with a payload containing a variable-length list of request objects.
2. The Experience API uses a For Each scope to split the list into individual objects and sends each object as a `REQU_OBJECT` message to an

Anypoint MQ queue named `REQUEST`.

3. The System API listens on the `REQUEST` queue, processes each message independently of all other messages, and sends a response message to a response queue named `RESPONSE_OBJECT`. The System API processing times can vary and are occasionally fairly long.
4. The Experience API listens on the `RESPONSE_OBJECT` queue and receives, transforms, and accumulates each response.
5. After receiving all the responses corresponding to the original `REQU` message, the Experience API creates and publishes a response Anypoint MQ message (`RESP`) to an Anypoint MQ queue named `RESPONSE` with a payload containing the list of responses returned by the System API, organized in the same order as the request objects originally sent in the `REQU` message.

Assume successful response messages are returned by Experience API to the `RESPONSE` queue for all `REQU` messages.

What is required so the Experience API can ensure the length and order of the list of objects in `RESP` and `REQU` match, while at the same time maximizing message throughput?

- ☒ A. Perform all communication involving the System API synchronously from within the For Each scope so that responses in `RESPONSE_OBJECT` are in the exact same order as request objects in the `RESP` message. Within the For Each scope, use an Async scope configured with a persistent object store.



Review



Comment

Back

Next

Choose 1 option.

occasionally fairly long.

4. The Experience API listens on the RESPONSE\_OBJECT queue and receives, transforms, and accumulates each response.

5. After receiving all the responses corresponding to the original REQU message, the Experience API creates and publishes a response Anypoint MQ message (RESP) to an Anypoint MQ queue named RESPONSE with a payload containing the list of responses returned by the System API, organized in the same order as the request objects originally sent in the REQU message.

Assume successful response messages are returned by Experience API to the RESPONSE queue for all REQU messages.

What is required so the Experience API can ensure the length and order of the list of objects in RESP and REQU match, while at the same time maximizing message throughput?

- ☐ A. Perform all communication involving the System API synchronously from within the For Each scope so that responses in RESPONSE\_OBJECT are in the exact same order as request objects in the RESP message. Within the For Each scope, use an Async scope configured with a persistent object store.
- ☐ B. Use a Scatter-Gather within the For Each scope to ensure response message order. Configure the Scatter-Gather with a persistent object store.
- ☒ C. After the For Each scope, add a second For Each scope configured with a persistent object store to collect response messages in the order in which they arrive, and then send the RESP message using this list of responses.
- ☐ D. Use a persistent object store to keep track of the list length and all object indexes in the REQU message, both in the For Each scope and in all communication involving the System API. Set the maxConcurrency parameter for the receiving flow in the Experience API to a value greater than 1.



Review



Comment

Back

Next

Choose 1 option.

A stock trading company handles millions of trades a day and requires excellent performance and reliability within its stock trading system. The company operates a number of event-driven APIs implemented as Mule applications that are hosted on various customer-hosted Mule clusters and needs to enable message exchanges between the APIs within their internal network using shared message queues.

What is an effective way to meet the cross-cluster messaging requirements of its event-driven APIs?

- ☐ A. Persistent VM queues with automatic acknowledgements
- ☒ B. eXtended Architecture (XA) transactions and XA connected components with manual acknowledgements
- ☐ C. JMS transactions with automatic acknowledgements
- ☐ D. Non-transactional JMS operations with a reliability pattern and manual acknowledgements



Review



Comment


Back

Next

Choose 1 option.

Refer to the exhibit.

In the flow of this Mule application, what thread pool handles requests to the HTTP Listener, and what thread pool handles the Logger component?



- ☐ A. BLOCKING\_IO, UBER
- ☒ B. Shared NIO Selector Pool, CPU\_LITE
- ☐ C. UBER, Dedicated NIO Selector Pool
- ☐ D. CPU\_LITE, CPU\_INTENSIVE



Review



Comment

Back

Next

Choose 1 option.

An organization plans to migrate all its Mule applications to Runtime Fabric (RTF). Currently, all Mule applications have been deployed to CloudHub using automated CI/CD scripts.

What steps should be taken to properly migrate the applications from CloudHub to RTF, while keeping the same automated CI/CD deployment strategy?

- ☐ A. A *runtimeFabric* dependency should be added as a mule-plugin to the pom.xml file in all the Mule applications.  
A *--runtimeFabric* command-line parameter should be added to the CI/CD deployment scripts.
- ☐ B. A *runtimeFabricDeployment* profile should be added to Mule configuration properties  
YAML files in all the Mule applications.  
CI/CD scripts must be modified to use the new configuration properties.
- ☒ C. A *runtimeFabricDeployment* profile should be added to the pom.xml file in all the Mule applications.  
CI/CD scripts must be modified to use the new RTF profile.
- ☐ D. The pom.xml and Mule configuration YAML files can remain unchanged in each Mule application.  
A *--runtimeFabric* command-line parameter should be added to the CI/CD deployment scripts



Review



Comment

Back

Next

Question 50 of 60

Choose 1 option.

A company is modernizing its legacy systems to accelerate access to applications and data while supporting the adoption of new technologies. The key to achieving this business goal is unlocking the company's systems and data, including a set of existing services hosted on-premises that can be accessed by authorized external clients. The IT staff is mainly experienced only with supporting its legacy systems.

Considering the current aggressive backlog and the project delivery requirements, the company wants to take a strategic approach in the first phase of its transformation projects by quickly deploying APIs in Mule runtimes that are able to scale, connect to on-premises systems, and migrate as needed.

Following MuleSoft best practices, what MuleSoft runtime deployment option best meets the company's goals to begin its digital transformation journey?

- ☐ A. CloudHub runtimes
- ☒ B. Customer-hosted runtimes provisioned by a MuleSoft services partner
- ☐ C. Runtime Fabric on VMs/bare metal
- ☐ D. Customer-hosted self-provisioned runtimes



Review



Comment

Back

Next

Choose 1 option.

An organization plans to migrate its deployment environment from an on-premises cluster to a Runtime Fabric (RTF) cluster. The on-premises Mule applications are currently configured with persistent object stores.

There is a requirement to enable Mule applications deployed to the RTF cluster to store and share data across application replicas and through restarts of the entire RTF cluster.

How can these reliability requirements be met?

- ☐ A. Replace persistent object stores with persistent VM queues in each Mule application deployment
- ☐ B. Install the Object Store pod on one of the cluster nodes
- ☒ C. Configure the Persistence Gateway in the RTF installation
- ☐ D. Configure Anypoint Object Store v2 to share data between replicas in the RTF cluster



Review



Comment

Back

Next

Choose 1 option.

How does a timeout attribute in a configured transaction manager help inform design decisions while using a JMS connector listening for incoming messages in an eXtended Architecture (XA) transaction managed by the transaction manager?

- ☒ A. The timeout defines the time that is allowed to pass without the transaction ending explicitly, and after the timeout expires, the transaction rolls back
- ☐ B. After the timeout is exceeded, stale JMS consumer threads are destroyed and new threads are created
- ☐ C. The timeout specifies the time allowed to pass between receiving JMS messages on the same JMS connection, and then after the timeout, a new JMS connection is established
- ☐ D. The time allowed to pass between committing the transaction and the completion of the Mule flow, and then after the timeout, flow processing triggers an error



Review



Comment

Back


Next

Choose 1 option.

Refer to the exhibit.

A customer is running Mule applications on Runtime Fabric for Self-Managed Kubernetes (RTF-BYOK8) in a multi-cloud environment.

Based on this configuration, how do Agents and Runtime Manager communicate, and what is exchanged between them?



Review




Comment

Back

Next

Choose 1 option.



- ☐ A. Runtime Manager initiates traffic to Agents to check the status of clusters and auto-deploys new images when updates to Mule applications are available
- ☐ B. Runtime Manager communicates with Agents to provide updates on cluster status, download images, and update Mule applications (when available and chosen to do so by customer) from the container registry
- ☒ C. Agents communicate with Runtime Manager to provide updates on cluster status, download images, and update Mule applications (automatically when available) from the container registry
- ☐ D. Agents communicate with Runtime Manager to provide updates on cluster status, download images, and update Mule applications (when available and chosen to do so by customer) from the container registry



Review



Comment

Back

Next

Choose 1 option.

An organization is struggling with frequent plugin version upgrades and external plugin project dependencies. The team wants to minimize the impact on applications by creating best practices that will define a set of default dependencies across all new and in-progress projects.

How can these best practices be achieved with the applications having the least amount of responsibility?

- ☐ A. Create a Mule domain project with all the dependencies defined in its pom.xml file and add each application to the domain project
- ☒ B. Create a parent POM with all the required dependencies and reference it in each application's pom.xml file
- ☐ C. Add all the dependencies in each application's pom.xml file
- ☐ D. Create a Mule plugin project with all the dependencies and add it as a dependency in each application's pom.xml file



Review



Comment

Back

Next

Choose 1 option.

A cluster with two customer-hosted Mule runtimes is hosting an application that has a flow with a JMS listener configured to subscribe to messages from a topic destination.

How should the JMS listener's `primaryNodeOnly` parameter and subscription parameter be configured so the JMS listener can receive messages in all the nodes of the cluster but only process each message once?

- ☒ A. Use the parameter `primaryNodeOnly="false"` on the JMS listener with a shared subscription
- ☐ B. Use the parameter `primaryNodeOnly="true"` on the JMS listener with a non-shared subscription
- ☐ C. Use the parameter `primaryNodeOnly="true"` on the JMS listener with a shared subscription
- ☐ D. Use the parameter `primaryNodeOnly="false"` on the JMS listener with a non-shared subscription



Review



Comment

Back

Next

### Salesforce Certified MuleSoft Platform Integration Architect Score Report

**CANDIDATE:** Frazier Senbetu

**DATE:** 10/22/2025

**CANDIDATE ID:** a9WRf000000scEnMAI

**REGISTRATION NUMBER:** 519095508

**EXAM CODE:** Mule-Arch-202

**EXAM:** Salesforce Certified MuleSoft Platform Integration Architect

**GRADE:** Pass

The report below shows the percentage of questions in each section that you answered correctly for the Salesforce Certified MuleSoft Platform Integration Architect exam. You will receive this information in your exam results email and can also access it via My Dashboard in Trailhead Academy.

**SECTION NAME:**

**PERCENT CORRECT:**

|                                                                  |      |
|------------------------------------------------------------------|------|
| Initiating integration Solutions on the Anypoint Platform        | 80%  |
| Designing for the Runtime Plane Technology Architecture          | 89%  |
| Designing Architecture Using Integration Paradigms               | 100% |
| Designing and Developing Mule Applications                       | 56%  |
| Designing Automated Tests for Mule Applications                  | 100% |
| Designing Integration Solutions to Meet Persistence Requirements | 83%  |
| Designing Integration Solutions to Meet Reliability Requirements | 80%  |
| Designing Integration Solutions to Meet Performance Requirements | 100% |
| Designing Integration Solutions to Meet Security Requirements    | 60%  |
| Applying DevOps Practices and Operating Integration Solutions    | 88%  |



Close