

Salesforce Certified MuleSoft Platform Architect

Question 1

When can CloudHub Object Store v2 be used?

- A. To store an unlimited number of key-value pairs
- B. To store key-value pairs with keys up to 300 characters
- C. To store payloads with an average size greater than 15MB
- D. To store information in Mule 4 Object Store v1

Question 2

A company deployed an API to a single worker/replica in the shared cloud in the U.S. West Region.

What happens when the Availability Zone experiences an outage?

- A. CloudHub will auto-redeploy the API in another Availability Zone in the U.S. West Region
- B. The Anypoint Platform admin is alerted when the API is experiencing an outage and needs to trigger the C17CD pipeline to redeploy to the U.S. East Region
- C. The API will be unavailable until the availability comes back online, at which time the worker/replica will be auto-restarted
- D. CloudHub will auto-redeploy the API in the U.S. East Region

Question 3

1 option.

- Route traffic via Dedicated load balancer (DLBs)
- Separate environments into production and non-production

Which solution meets the customer's needs?

- A. One production and one non-production Virtual Private Cloud (VPC) per Business group. Minimize CIDR aligning with projected application total. Divide availability zones during deployment of APIs for HA.
- B. One production and one non-production Virtual Private Cloud (VPC). Configure subnet to differentiate between business groups. Allocate maximum CIDR per VPCs to make it easier to add Child groups. Span VPC to cover three availability zones.
- C. One production and one non-production Virtual Private Cloud (VPC). Use availability zones to differentiate between Business groups. Allocate maximum CIDR per VPCs to ensure HA across availability zones

D. One production and one non-production Virtual Private Cloud (VPC) per Business group. Minimize CIDR aligning with projected application total. Choose a MuleSoft CloudHub 1.0 region with multiple availability zones. Deploy multiple workers for HA.

Question 4

An established communications company is beginning its API-led connectivity journey. The company has been using a successful Enterprise Data Model for many years. The company has identified a self-service account management app as the first effort for API-led, and it has identified the following APIs.

Experience layer: Mobile Account Management EAPI, Browser Account Management EAPI

Process layer: Customer Lookup PAPI, Service Lookup PAPI, Account Lookup PAPI

System layer: Customer SAPI, Account SAPI, Product SAPI, Service SAPI

According to MuleSoft's API-led connectivity approach, which API would not be served by the Enterprise Data Model?

A. Mobile Account Management EAPI

B. Service SAPI

C. Customer SAPI

D. Customer Lookup PAPI

Question 5

Which statement is true about where an API policy is defined in Anypoint Platform and how it is applied to API instances?

A. The API policy is defined in API Manager and then automatically applied to all API instances

B. The API policy is defined in Runtime Manager as part of the API deployment to a Mule runtime and then only applied to the specific API instance

C. The API policy is defined in API Manager for a specific API instance and then only applied to the specific API instance

D. The API policy is defined in API Manager and then applied to all API instances at deployment time

Question 6

A circuit breaker strategy is planned in order to meet the goal of improved response time and demand on a downstream API.

- Circuit Open: More than 10 errors per minute for three minutes

- Circuit Half-Open: One error per minute

- Circuit Closed: Less than one error per minute for five minutes

Out of several proposals from the engineering team, which option will meet this goal?

A. Implement the strategy in a Mule application, and provide the settings in the YAML configuration

- B. Add the Circuit Breaker policy to the API instance, and configure the required settings
- C. Create Anypoint Monitoring alerts for Circuit Open/Closed configurations, and then implement a retry strategy for Circuit Half-Open configuration
- D. Create a custom policy that implements the circuit breaker and includes policy template expressions for the required settings

Question 7

A developer for a transportation organization is implementing exactly one processing functionality in a Reservation Mule application to process and store passenger records. This Reservation application will be deployed to multiple CloudHub workers/replicas. It is possible that several external systems could send duplicate passenger records to the Reservation application.

An appropriate storage mechanism must be selected to help the Reservation application process each passenger record exactly once as much as possible. The selected storage mechanism must be shared by all the CloudHub workers/replicas in order to synchronize the state information to assist attempting exactly once processing of each passenger record by the deployed Reservation Mule application.

Which type of simple storage mechanism in Anypoint Platform allows the Reservation Mule application to update and share data between the CloudHub workers/replicas exactly once, with minimal development effort?

- A. Runtime Fabric Object Store
- B. In-memory Mule Object Store
- C. Non-persistent Object Store
- D. Persistent Object Store

Question 8

An API implementation is deployed to CloudHub.

What conditions can be alerted on using the default Anypoint Platform functionality, where the alert conditions depend on the API invocations to an API implementation?

- A. When the API invocations are not over a secure TLS/SSL communication channel
- B. When the API invocations originate from a geography different than the API
- C. When the number of API invocations are below a threshold
- D. When the API invocations are sent directly to the internal DNS record of the API implementation

Question 9

A developer from the Central IT team has created an initial version of the RAML definition in Design Center for an OAuth 2.0-protected System API and published it to Exchange. Another developer from LoB IT discovered the System API in Exchange and would like to leverage it in the Process API.

What is the MuleSoft-recommended approach for Process API to invoke the System API?

A. The Process API uses the REST Connect Connector autogenerated in Exchange for the System API

B. The Process API uses property YAML files to store the System API URLs and uses the HTTP Request Connector to invoke the System API

C. The Process API needs to import an OAuth 2.0 module from Exchange first and update it with OAuth 2.0 credentials before the System API can be invoked

D. The Process API manually updates the Process API POM file to include the System API as a dependency

Question 10

A customer has an ELA contract with MuleSoft. An API deployed to CloudHub is consistently experiencing performance issues. Based on the root cause analysis, it is determined that autoscaling needs to be applied.

How can this be achieved?

A. Configure a policy so that when the response time reaches a certain threshold the worker/replica type increases (vertical scaling)

B. Configure two separate policies: When CPU and memory reach a certain threshold, increase the worker/replica type (vertical scaling) and the number of workers/replicas (horizontal scaling)

C. Configure a policy based on CPU usage so that CloudHub auto-adjusts the number of workers/replicas (horizontal scaling)

D. Configure a policy so that when the number of HTTP requests reaches a certain threshold the number of workers/replicas increases (horizontal scaling)

Question 11

An application updates an inventory running only one process at any given time to keep the inventory consistent. This process takes 200 milliseconds (.2 seconds) to execute; therefore, the scalability threshold of the application is five requests per second.

What is the impact on the application if horizontal scaling is applied, thereby increasing the number of Mule workers?

A. The application scalability threshold is now 10 requests per second

B. The application scalability threshold is five requests per second regardless of the horizontal scaling

C. The total process execution time is now 100 milliseconds (.1 seconds)

D. Horizontal scaling cannot be applied to an already-running application

Question 12

A Mule application implements an API. The Mule application has an HTTP Listener whose connector configuration sets the HTTPS protocol and hard-codes the port value. The Mule application is deployed to an Anypoint VPC and uses the CloudHub 1.0 Shared Load Balancer (SLB) for all incoming traffic.

Which port number must be assigned to the HTTP Listener's connector configuration so that the Mule application properly receives HTTPS API invocations routed through the SLB?

A. 8082

B. 80

C. 8092

D. 443

Question 13

An enterprise is embarking on the API-led digital transformation journey, and the central IT team has started to define System APIs. Currently there is no Enterprise Data Model being defined within the enterprise, and the definition of a clean Bounded Context Data Model requires too much effort.

According to MuleSoft's recommended guidelines, how should the System API data model be defined?

A. The System APIs should expose all back-end system fields

B. The data model of the System APIs should make use of data types that approximately mirror those from the back-end systems

C. If there are misspellings of the data fields in the back-end system, System APIs should not correct it, and expose it as-is to mirror the back-end systems

D. The data model should define its own naming convention, and not follow the same naming as the back-end systems

Question 14

A customer wants to monitor and gain insights about the number of requests coming in a given time period as well as to measure key performance indicators (response times, CPU utilization, number of active APIs).

Which tool provides these data insights?

A. Anypoint Monitoring

B. Runtime Alerts

C. Functional Monitoring

D. API Manager

Question 15

An organization has built an application network following the API-led connectivity approach recommended by MuleSoft. To protect the application network against attacks from malicious external API clients, the organization plans to apply JSON Threat Protection policies.

To which API-led connectivity layer should the JSON Threat Protection policies most commonly be applied?

A. Process layer

B. Experience layer

C. All layers

D. System layer

Question 16

Which three tools automate the deployment of Mule applications?

A. API Community Manager

B. Anypoint Studio

C. Platform APIs

D. Anypoint Platform CLI

E. Runtime Manager

F. Mule Maven plugin

Question 17

A Production environment is running on a dedicated Virtual Private Cloud (VPC) on CloudHub 1.0, and the security team guidelines clearly state no traffic on HTTP.

Which two options support these security guidelines?

A. Remove the entry from the VPC firewall rule

```
{  
  "CIDR Block": "0.0.0.0/0"; // (Anywhere)  
  "Protocol": "TCP";  
  "From port": 8081,  
}  
{  
  "CIDR Block": "10.111.0.0/24"; // (Local VPC)  
  "Protocol": "TCP";  
  "From port": 8081,  
}
```

B. Configure the IP Blocklist policy to control access to a configured API endpoint from either a single IP address or a range of IP addresses

Question 18

An organization requires several APIs to be secured with OAuth 2.0, and PingFederate has been identified as the identity provider for API client authorization. The PingFederate Client Provider is configured in access management, and the PingFederate OAuth 2.0 Token Enforcement policy is configured for the API instances required by the organization. The API instances reside in two business groups (Group A and Group B) within the Master Organization (Master Org).

What should be done to allow API consumers to access the API instances?

A. The API consumer should create a client application and request access to the API in Anypoint Exchange, and the API administrator should approve the request

B. The API consumer should create a client application and request access to the API in Ping Identity, and the organization's Ping Identity workflow will grant access

C. The API administrator should grant access to the API consumers by creating contracts in the relevant API instances in API Manager

D. The API administrator should configure the correct client discovery URL in both child business groups, and the API consumer should request access to the API in Ping Identity

Question 19

Which statement is true about Spike Control policy and Rate Limiting policy?

A. All requests are rejected after the limit is reached in Rate Limiting policy, whereas the requests are queued in Spike Control policy after the limit is reached

B. To protect Experience APIs by limiting resource consumption, Rate Limiting policy must be applied

C. In a clustered environment, the Rate Limiting and Spike Control policies are applied to each node in the cluster

D. In order to apply Rate Limiting and Spike Control policies, a contract to bind client application and API is needed for both

Question 20

A large lending company has developed an API to unlock data from a database server and web server. The API has been deployed to Anypoint Virtual Private Cloud (VPC) on CloudHub 1.0.

The database server and web server are in the customer's secure network and are not accessible through the public internet. The database server is in the customer's AWS VPC, whereas the web server is in the customer's on-premises corporate data center.

How can access be enabled for the API to connect with the database server and the web server?

A. Set up VPC peering with the customer's on-premises corporate data center and a VPN tunnel to AWS VPC

B. Set up VPC peering with AWS VPC and a VPN tunnel to the customer's on-premises corporate data center

C. Set up VPC peering with AWS VPC and the customer's on-premises corporate data center

D. Set up a transit gateway to the customer's on-premises corporate data center through AWS VPC

Question 21

Which statement is true about identity management and client management on Anypoint Platform?

- A. If an external identity provider is configured, the SAML 2.0 bearer tokens issued by the identity provider cannot be used for invocations of the Anypoint Platform web APIs
- B. If an external client provider is configured, it must be configured at the Anypoint Platform organization level and cannot be assigned to individual business groups and environments
- C. Anypoint Platform supports configuring one external identity provider
- D. Both client management and identity management require an identity provider

Question 22

An online store's marketing team has noticed an increase in customers leaving online baskets without checking out. They suspect a technology issue is at the root cause of the baskets being left behind. They approach the Center for Enablement to ask for help identifying the issue. Multiple APIs from across all the layers of their application network are involved in the shopping application.

Which feature of the Anypoint Platform can be used to view metrics from all involved APIs at the same time?

- A. API Manager
- B. Custom dashboards
- C. Built-in dashboards
- D. Functional monitoring

Question 23

Refer to the exhibits. Which architectural constraint is compatible with the API-led connectivity architectural style?

- System API(s)
- Business process
- Process API(s)
- Experience API(s)
- API Consumer

A. Use a Process API to orchestrate calls to multiple System APIs but not to other Process APIs

- API Experience
- Process API(s)

Question 24

A customer wants to host their MuleSoft applications in CloudHub 1.0, and these applications should be available at the domain <https://api.acmecorp.com>.

After creating a dedicated load balancer (DLB) called acme-dlb-prod, which further action must the customer take to complete the configuration?

- A. Configure the DLB with a TLS certificate for api.acmecorp.com and create a CNAME record from api.acmecorp.com to acme-dlb-prod.e1.cloudhub.io
- B. Configure the DLB with a TLS certificate for api.acmecorp.com and create an A record for api.acmecorp.com to the public IP addresses associated with their DLB
- C. Configure the DLB with a TLS certificate for acme-dlb-prod.lb.anypointdns.net and create a CNAME record from api.acmecorp.com to acme-dlb-prod.lb.anypointdns.net
- D. Configure the DLB with a TLS certificate for api.acmecorp.com and create a CNAME record from api.acmecorp.com to acme-dlb-prod.lb.anypointdns.net

Question 25

Refer to the exhibits.

The Line of Business (LoB) of an eCommerce company is requesting a process that sends automated notifications via email every time a new order is processed through the customer's mobile application or through the internal company's web application. In the future, multiple notification channels may be added: for example, text messages and push notifications.

What is the most effective API-led connectivity approach for the scenario described above?

- A. Create one Experience API for the web application and one for the mobile application.

Create a Process API to orchestrate, retrieve the email template from a database, and send the email using the Anypoint Connector for Email.

Experience APIs

- Process APIs
- Mobile API
- Web app API

Order Notification

Question 26

A Platform Architect inherits a legacy monolithic SOAP-based web service that performs a number of tasks, including showing all policies belonging to a client. The service connects to two back-end systems – a life-insurance administration system and a general insurance administration system – and then queries for insurance policy information within each system, aggregates the results, and presents a SOAP-based response to a user interface (UI).

The architect wants to break up the monolithic web service to follow API-led conventions.

Which part of the service should be put into the process layer?

- A. Combining the insurance policy information from the administration systems

- B. Authenticating and maintaining connections to each of the back-end administration systems
- C. Querying the data from the administration systems
- D. Presenting the SOAP-based response to the UI

Question 27

An API with multiple API implementations (Mule applications) is deployed to both CloudHub and customer-hosted Mule runtimes. All the deployments are managed by the Mulesoft-hosted control plane. An alert needs to be triggered whenever an API implementation stops responding to API requests, even if no API clients have called the API implementation for some time.

What is the most effective out-of-the-box solution to create these alerts to monitor the API implementations?

- A. Create monitors in Anypoint Functional Monitoring for the API implementations, where each monitor repeatedly invokes an API implementation endpoint
- B. Handle API invocation exceptions within the calling API client and raise an alert from that API client when such an exception is thrown
- C. Add code to each API client to send an Anypoint Platform REST API request to generate a custom alert in Anypoint Platform when an API invocation times out
- D. Configure one Worker Not Responding alert in Anypoint Runtime Manager for all API implementations that will then monitor every API implementation

Question 28

An eCommerce company is adding a new Product Details feature to their website. A customer will launch the product catalog page, a new Product Details link will appear by product where they can click to retrieve the product detail description. Product details are updated with product update releases, once or twice a year. Presently the database response time has been very slow due to high volume.

What action retrieves the product details with the lowest response time, fault tolerant, and consistent data?

- A. Select the product details from a database in a Cache scope and return them within the API response
- B. Select the product details from a database and return them within the API response
- C. Select the product details from a database and put them in Anypoint MQ; the Anypoint MQ subscriber will receive the product details and return them within the API response
- D. Use an object store to store and retrieve the product details originally read from a database and return them within the API response

Question 29

Which component monitors APIs and endpoints at scheduled intervals, receives reports about whether tests pass or fail, and displays statistics about API and endpoint performance?

A. API Functional Monitoring

- B. API Analytics
- C. Anypoint Monitoring dashboards
- D. Anypoint Runtime Manager alerts

Question 30

Which out-of-the-box key performance indicator measures the success of a typical Center for Enablement and is immediately available in responses from Anypoint Platform APIs?

- A. Per deployed API implementation, the amount of bandwidth consumed each day
- B. Per published API, the number of consumers that requested access to the API and have been approved in the Production environment**
- C. Per business group, the ratio of the number of production API implementations deployed using a CI/CD pipeline to the number of production API implementations deployed manually
- D. Per published API, the number of developers that downloaded a version of the API specification

Question 31

Which scenario is suited for MUnit tests instead of integration tests?

- A. For read-only interactions to any dependencies (such as other web APIs)**
- B. When no mocking is permissible
- C. For tests that are implemented using SoapUI
- D. When testing does not require knowledge of implementation details

Question 32

A large organization with an experienced central IT department is getting started using MuleSoft. There is a project to connect a siloed back-end system to a new Customer Relationship Management (CRM) system. The Center for Enablement is coaching them to use API-led connectivity.

What action would support the creation of an application network using API-led connectivity?

- A. To expedite this project, central IT should extend the CRM system and back-end systems to connect to one another using built in integration interfaces
- B. Determine if the new CRM system supports the creation of custom REST APIs, establishes a private network with CloudHub, and supports OAuth 2.0 authentication
- C. Invite the business analyst to create a business process model to specify the canonical data model between the two systems
- D. Create a System API to unlock the data on the back-end system using a REST API**

Question 33

A large company wants to implement IT infrastructure in its own data center, based on the corporate IT policy requirements that data and metadata reside locally.

Which combination of Mule control plane and Mule runtime plane(s) meets the requirements?

- A. The MuleSoft-hosted control plane and Anypoint Runtime Fabric for the runtime plane
- B. Anypoint Platform Private Cloud Edition for the control plane and customer-hosted Mule runtimes for the runtime plane
- C. The MuleSoft-hosted control plane and customer-hosted Mule runtimes for the runtime plane
- D. Anypoint Platform Private Cloud Edition for the control plane and the MuleSoft-hosted runtime plane

Question 35

When should idempotency be taken into account?

- A. When storing the results of a previous request for use in response to subsequent requests
- B. When preventing duplicate processing from multiple sent requests
- C. When sending concurrent update requests for the same entity
- D. When making requests to update currently locked entities

Question 36

An API is protected with a Client ID Enforcement policy and uses the default configuration. Access is requested for the client application to the API, and an approved contract now exists between the client application and the API.

How can a consumer of this API avoid a 401 error "Unauthorized or invalid client application credentials"?

- A. Send the obtained token as a header in every call
- B. Send the obtained client_id and client_secret in the request body
- C. Send the obtained client_id and client_secret in the header of every API Request call
- D. Send the obtained client_id and client_secret as URI parameters in every call

Question 37

Which two statements are true about the technology architecture of an Anypoint Virtual Private Cloud (VPC)?

- A. CIDR blocks are used
- B. Round-robin load balancing is used to distribute client requests across different applications
- C. By default, HTTP requests can be made from the public internet to workers at port 8091
- D. Ports 8081 and 8082 are used

E. Anypoint VPC is responsible for load balancing the applications

Question 38

A company deploys Mule applications with default configurations through Runtime Manager to customer-hosted Mule runtimes. Each Mule application is an API implementation that exposes RESTful interfaces to API clients. The Mule runtimes are managed by the MuleSoft-hosted control plane. The payload is never used by any Logger components.

When an API client sends an HTTP request to a customer-hosted Mule application, which metadata or data (payload) is pushed to the MuleSoft-hosted control plane?

A. Only the metadata

B. The data and metadata

C. No data

D. Only the data

Question 39

A company is using an on-prem cluster in the data center as a runtime plane and MuleSoft-hosted control plane.

How can the company monitor the detailed performance metrics on the Mule applications deployed to the cluster from the control plane?

A. Due to the potential performance impact on the runtime nodes, the Monitoring agent should be installed on a separate server

B. There is no action needed as the on-prem runtime automatically sends the performance data to the control plane

C. Monitoring Agent must be installed on each node in the cluster

D. The settings of the Monitoring section in the control plane must be updated to enable detailed logging on the metrics to be captured

Based on the images provided, here is the complete question and its available options.

Question 40

A system has two legacy backends (Legacy System-1 and Legacy System-2) with their own System APIs (SAPI-1, SAPI-2). There is an existing Process API (PAPI-1) that orchestrates these SAPIs. The goal is to introduce a new Process API (PAPI-2) and migrate clients from the old PAPI-1 to the new PAPI-2.

****A.**** Deploy an API Proxy with an endpoint from API Manager that redirects the existing PAPI endpoints to the new PAPI.

****B.**** Refactor the PAPI implementation into an Experience API (EAPI). Add data transformation and routing logic to the EAPI implementation to support the required endpoints of both mobile and web clients. Postpone deployment until all API consumers acknowledge they are ready to migrate to the new API version.

Question 41

Refer to the exhibit.

Which API policy would least likely be used when designing an Experience API that is intended to work with a consumer mobile phone or tablet application?

- invokes
- Mobile app(s)
- Experience API
- Least likely policy?

A. Client ID Enforcement

B. JSON Threat Protection

C. IP Allowlist

D. OAuth 2.0 Access Token Enforcement

Question 42

Refer to the exhibit.

A Mule 4 API has been deployed to CloudHub and a Basic Authentication - Simple policy has been applied to all API methods and resources. However, the API is still accessible by clients without using authentication.

How is this possible?

Apply Basic authentication - Simple policy

Enforces HTTP Basic authentication according to the details configured in the policy

User Name *

User Password *

Method & Resource conditions

- Apply configurations to all API methods & resources

- Apply configurations to specific methods & resources

Question 43

A Platinum customer uses the U.S. control plane and deploys applications to CloudHub in Singapore with a default log configuration.

The compliance officer asks where the logs and monitoring data reside?

- A. Logs and monitoring data are held in the United States
- B. Logs and monitoring data are held in Singapore
- C. Logs are held in the United States and monitoring data is held in Singapore
- D. Logs are held in Singapore and monitoring data is held in the United States

Question 44

--- should be on the agenda of the first meeting?

- A. Define C4E objectives, mission statement, guiding principles, and responsibilities
- B. Specify operating model for the MuleSoft Integrations division
- C. A walk through of common-services best practices for logging, auditing, exception handling, caching, security via policy, and rate limiting/throttling via policy
- D. Explore API monetization options based on identified use cases through MuleSoft

Question 45

A company is building an application network using MuleSoft's recommendations for various API layers.

What is the main (default) role of a process API in an application network?

- A. To manage and process the secure direct communication between a back-end system and an end-user client or mobile device in the application network
- B. To secure and optimize the data synchronization processing of large data dumps between back-end systems
- C. To automate parts of business processes by coordinating and orchestrating the invocation of other APIs in the application network
- D. To secure, manage, and process communication with specific types of end-user client applications or devices in the application network

Question 46

A new upstream API is being designed to offer a service-level agreement (SLA) of 500ms median and 800ms maximum (99th percentile) response time. The corresponding API implementation must sequentially invoke three downstream APIs of equivalent complexity.

The first of these downstream APIs offers the following SLA for its response time. Median: 100ms; 80th percentile: 500ms; 95th percentile: 1,000ms.

If possible, how can a timeout be set in the upstream API for the invocation of the first downstream API to meet the new upstream API's desired SLA; if this cannot be done, what should be done to best support the new upstream API?

Proposed SLA:

500 ms median and 800 ms maximum (99th percentile)

Downstream SLA:

median: 100 ms, 80th percentile: 500 ms, 95th percentile: 1000 ms

API Consumer

- Unstream
- Downstream API 1
- Downstream API 2
- Downstream API 3

Question 46 of 60

Choose 1 option.

- API Consumer
- Downstream API 1
- Upstream API
- Downstream API 2
- What timeout is needed?
- Downstream API 3

A. Set a timeout of 100ms; this leaves 400ms for the other two downstream APIs to complete

B. Set a timeout of 50ms; this times out more invocations of that API but gives additional room for retries

C. Do not set a timeout; the invocation of this API is mandatory and requires waiting until it responds

D. No timeout is possible to meet the upstream API's desired SLA; a different SLA must be negotiated with the first downstream API or an alternative API must be invoked

Question 47

An auto manufacturer has a mature CI/CD practice and wants to automate packaging and deployment of any Mule applications to various deployment targets, including CloudHub workers/replicas, customer-hosted Mule runtimes, and Anypoint Runtime Fabric.

Which MuleSoft-provided tool or component facilitates automating the packaging and deployment of Mule applications to various deployment targets as part of the company's CI/CD practice?

- A. Anypoint Runtime Manager
- B. Anypoint Platform CLI
- C. Anypoint Platform REST APIs
- D. Mule Maven plugin**

****Question 48**

Several times a week, an API implementation shows several thousand requests per minute in an Anypoint Monitoring dashboard. Between these bursts, the dashboard shows between two and five requests per minute. The API implementation is running on Anypoint Runtime Fabric with two non-clustered replicas, reserved vCPU 1.0 and vCPU Limit 2.0.

An API consumer has complained about slow response time, and the dashboard shows the 99 percentile is greater than 120 seconds at the time of the complaint. It also shows greater than 90% CPU usage during these time periods.

In manual tests in the QA environment, the API consumer has consistently reproduced the slow response time and high CPU usage, and there were no other API requests at this time. In a brainstorming session, the engineering team has created several proposals to reduce the response time for requests.

Which proposal should be pursued first?

- A. Modify the API client to split the problematic request into smaller, less-demanding requests
- B. Throttle the API client to reduce the number of requests per minute
- C. Increase the number of replicas of the API implementation**
- D. Increase the vCPU resources of the API implementation

****Question 49**

A Rate Limiting policy is applied to an API implementation to protect the back-end system. Recently, there have been surges in demand that cause some API client POST requests to the API implementation to be rejected with policy-related errors, causing delays and complications to the API clients.

How should the API policies that are applied to the API implementation be changed to reduce the frequency of errors returned to API clients, while still protecting the back-end system?

- A. Keep the Rate Limiting policy and add an SLA-based Spike Control policy
- B. Remove the Rate Limiting policy and add a Spike Control policy**
- C. Keep the Rate Limiting policy and add a Client ID Enforcement policy
- D. Remove the Rate Limiting policy and add an HTTP Caching policy

****Question 50**

An existing Quoting API is defined in RAML and used by REST clients for interacting with the quoting engine. Currently there is a resource defined in the RAML that allows the creation of quotes; however, a new requirement was just received to allow for the updating of existing quotes.

Which two actions need to be taken to facilitate this change so it can be processed?

- A. Deprecate existing versions of the API in Exchange
- B. Update the RAML with new method details for the update request**
- C. Add a new API policy to API Manager to allow access to the updated endpoint
- D. Remove the old client applications and create new client applications to account for the changes
- E. Update the API implementation to accommodate the new update request**

****Question 51**

An IT Security Compliance Auditor is assessing which nonfunctional requirements (NFRs) are already being implemented to meet security measures.

- The Web API has Rate-Limiting SLA
- Basic Authentication - LDAP
- JSON Threat Protection
- IP Allowlist policies applied

Which two NFRs are enforced?

- A. The API is protected against XML invocation attacks
- B. Sensitive data is masked to prevent compromising critical information
- C. The API invocations are coming from a known subnet range**
- D. Username/password supported to validate login credentials**
- E. Performance expectations are to be allowed up to 1,000 requests per second

****Question 52**

A business process is being implemented within an organization's application network. The architecture group proposes using a more coarse grained application network design with relatively fewer APIs deployed to the application network compared to a more fine grained design.

Overall, which factor typically increases with a more coarse grained design for this business process implementation and deployment compared with using a more fine-grained design?

- A. The complexity of each API implementation**
- B. The usage of network infrastructure resources by the application network
- C. The number of discoverable assets related to APIs deployed in the application network
- D. The number of possible connections between API implementations in the application network

****Question 53**

Refer to the exhibit.

An Order API triggers a sequence of other API calls to look up details of an order's items in a back-end inventory database. The Order API calls the OrderItems process API, which calls the Inventory system API. The Inventory system API performs database operations in the back-end inventory database.

The network connection between the Inventory system API and the database is known to be unreliable and hang at unpredictable times.

Where should a two-second timeout be configured in the API processing sequence so that the Order API never waits more than two seconds for a response from the OrderItems process API?

- Order API
- OrderItems process API
- Inventory system API
- Database

A. In the OrderItems process API implementation

B. In the Inventory system API implementation

****Question 54**

A TemperatureSensors API instance is defined in API Manager in the PROD environment of the CAR_FACTORY business group. An AcmeTemperatureSensors Mule application implements this API instance and is deployed from Runtime Manager to the PROD environment of the CAR_FACTORY business group. A policy that requires a valid client ID and client secret is applied in API Manager to the API instance.

Where can an API consumer obtain a valid client ID and client secret to call the AcmeTemperatureSensors Mule application?

- A. In API Manager, from the PROD environment of the CAR_FACTORY business group
- B. In access management, from the PROD environment of the CAR_FACTORY business group
- C. In Anypoint Exchange, from an API client application that has been approved for the TemperatureSensors API instance

D. In secrets manager, request access to the Shared Secret static username/password

****Question 55**

The asset version 2.0.0 of the Order API is successfully published in Exchange and configured in API Manager with the Autodiscovery API ID correctly linked to the API implementation. A new GET method is added to the existing API specification, and after updates, the asset version of the Order API is 2.0.1.

What happens to the Autodiscovery API ID when the new asset version is updated in API Manager?

- A. The API ID changes, so the API implementation must be updated with the latest API ID for the new asset version in the API Autodiscovery global element
- B. The API ID does not change, but the API implementation must be updated in the API Autodiscovery global element to indicate the new asset version 2.0.1
- C. The API ID changes, but no changes are needed to the API implementation for the new asset version in the API Autodiscovery global element because the API ID is automatically updated
- D. The API ID does not change, so no changes to the API implementation are needed for the new asset version in the API Autodiscovery global element

****Question 56**

A European company has customers all across Europe, and the IT department is migrating from an older platform to MuleSoft. The main requirements are that the new platform should allow redeployments with zero downtime and deployment of applications to multiple runtime versions, provide security and speed, and utilize Anypoint MQ as the message service.

Which runtime plane should the company select based on the requirements without additional network configuration?

- A. Anypoint Runtime Fabric on Self-Managed Kubernetes for the runtime plane
- B. Customer-hosted runtime plane
- C. MuleSoft-hosted runtime plane (CloudHub)
- D. Runtime Fabric on VMs / Bare Metal for the runtime plane

****Question 57**

A team is planning to enhance an Experience API specification, and they are following API-led connectivity design principles.

What is their motivation for enhancing the API?

- A. The underlying System API is updated to provide more detailed data for several heavily used resources
- B. An IP Allowlist policy is being added to the API instances in the Development and Staging environments
- C. A Canonical Data Model is being adopted that impacts several types of data included in the API
- D. The primary API consumer wants certain kinds of endpoints changed from the Center for Enablement standard to the consumer system standard

****Question 58**

A client has several applications running on the Salesforce service cloud. The business requirement for integration is to get daily data changes from Account and Case Objects. Data needs to be moved to the client's private cloud AWS DynamoDB Instance as a single JSON and the business foresees only wanting five attributes from the Account object, which has 219 attributes (some custom) and eight attributes from the Case Object.

What design should be used to support the API/ Application data model?

- A. Create separate entities for Account and Case Objects by mimicking all the attributes in SAPI, which are combined by the PAPI and filtered to provide JSON output containing 13 attributes.
- B. Start implementing an Enterprise Data Model by defining enterprise Account and Case Objects and implement SAPI and DynamoDB tables based on the Enterprise Data Model.
- C. Request client's AWS project team to replicate all the attributes and create Account and Case JSON table in DynamoDB. Then create separate entities for Account and Case Objects by mimicking all the attributes in SAPI to transfer JSON data to DynamoDB for respective Objects
- D. Create separate entities for Account with five attributes and Case with eight attributes in SAPI, which are combined by the PAPI to provide JSON output containing 13 attributes.

****Question 59**

A manufacturing company has deployed an API implementation to CloudHub and has not configured it to be automatically restarted by CloudHub when the worker is not responding.

Which statement is true when no API Client invokes that API implementation?

- A. Alerts on the API invocations are raised but no alerts on the API implementation can be raised
- B. Alerts on the API invocation and API implementation can be raised
- C. No alert on the API invocations is raised but alerts on the API implementation can be raised
- D. No alert on the API invocations and API implementation can be raised

****Question 60**

An operations team is analyzing the effort needed to set up monitoring of their application network. They are looking at which API invocation metrics can be used to identify and predict trouble without having to write custom scripts or install additional analytics software or tools.

Which type of metrics can satisfy this goal of directly identifying and predicting failures?

- A. The effectiveness of the application network based on the level of reuse
- B. The number and types of API policy violations per day
- C. The number and types of past API invocations across the application network
- D. The ROI from each API invocation

61. | I on | uses a | r-nc | customer-hosted Mule runtime cluster to host one stateless API implementation. The API is |

| ---|---|---|---|

| ces: | TTPS tt | igh: | ad balancer that uses round-robin for load distribution. Each node in the cluster has been sized |

| acc , | mes th: | irrei: | umber of requests. |

Two nodes in the cluster experience a power outage and are no longer available. The load balancer detects the outage and blocks the two unavailable nodes from receiving further HTTP requests.

What performance-related consequence is guaranteed to happen, on average, assuming the remaining cluster nodes are fully operational?

- A. 100% increase in the number of requests received by each remaining node
- B. 100% increase in the average response time of the API implementation
- C. 50% increase in the number of thread pools on the remaining servers
- D. 50% reduction in the throughput of the API implementation

62. An external API frequently invokes an Employees System API to fetch employee data from a MySQL database. The architect must design a caching strategy to query the database only when there is an update to the Employees table or else return a cached response in order to minimize the number of redundant transactions being handled by the database.

What must the architect do to achieve the caching objective?

- A. Use an On Table Row operation configured with the Employees table and call invalidate cache.

Use an object-store-caching-strategy and the default expiration interval.

- B. Use a Scheduler with a fixed frequency set to every hour, triggering an invalidate cache flow.

Use an object-store-caching-strategy and set the expiration interval to 1 hour.

- C. Use a Scheduler with a fixed frequency set to every hour to trigger an invalidate cache flow.

Use an object-store-caching-strategy and the default expiration interval.

- D. Use an On Table Row operation configured with the Employees table, call invalidate cache, and hardcode the new Employees data to cache.

Use an object-store-caching-strategy and set the expiration interval to 1 hour.

63. A large life sciences customer plans to use the Mule Tracing module with the Mapped Diagnostic Context (MDC) logging operations to enrich logging in its Mule application and to improve tracking by providing more context in the Mule application logs. The customer also wants to improve throughput and lower the message processing latency in its Mule application flows.

After installing the Mule Tracing module in the Mule application, how should logging be performed in flows in Mule applications, and what should be changed in the log4j2.xml files?

- A. In the flows, add Mule Tracing module Set logging variable operations before any Core Logger components. In log4j2.xml files, change the appender's pattern layout to use %MDC and then assign the appender to a Logger or Root element.
- B. In the flows, wrap Logger components in Async scopes. In log4j2.xml files, change the appender's pattern layout to use the %asyncLogger placeholder and then assign the appender to a Logger or Root element.
- C. In the flows, add Mule Tracing module Set logging variable operations before any Core Logger components. In log4j2.xml files, change the appender's pattern layout to use %asyncLogger placeholder and then assign the appender to an AsyncLogger element.
- D. In the flows, add Mule Tracing module Set logging variable operations before any Core Logger components. In log4j2.xml files, change the appender's pattern layout to use the %MDC placeholder and then assign the appender to an AsyncLogger element.

64. A retail company is implementing a MuleSoft API to get inventory details from two vendors by invoking each vendor's online applications. Due to network issues, the invocations to the vendor applications are timing out intermittently, but the requests are successful after re-invoking each vendor application.

What is the most performant way of implementing the API to invoke each vendor application and to retry invocations that generate timeout errors?

- A. Use a Scatter-Gather scope to invoke each vendor application on a separate route.
Use an Until-Successful scope in each route to retry requests that raise timeout errors.
- B. Use a Round-Robin scope to invoke each vendor application on a separate route.
Use a Try-Catch scope in each route to retry requests that raise timeout errors.
- C. Use a Choice scope to invoke each vendor application on a separate route.
Place the Choice scope inside an Until-Successful scope to retry requests that raise timeout errors.
- D. Use a For-Each scope to invoke the two vendor applications in series, one after the other.
Place the For-Each scope inside an Until-Successful scope to retry requests that raise timeout errors.

65. An insurance company is implementing a new Mule application that supports a set of critical functions driven by a REST-API-enabled claim payment rules engine hosted on Oracle ERP. The new Mule application requires many data transformation operations as it performs batch processing logic. The company wants to leverage and reuse as many of its existing Java-based capabilities (such as classes, objects, frameworks, and data models) as possible.

What approach should be used when implementing the required data mappings and transformations between other Mule application components and Oracle ERP?

- A. Transform by calling any suitable Java class from DataWeave

- B. Create new metadata RAML classes in the Mule runtime corresponding to existing Java objects, and then perform required transformations via DataWeave
- C. Invoke any of the appropriate Java methods directly, create corresponding metadata RAML classes, and then perform required transformations via DataWeave
- D. Transform via the XSLT module

66. What requirement prevents using Anypoint MQ as the messaging broker for a Mule application?

- A. When the messaging broker must be deployed on-premises

- B. When the payload sent through the message broker must be encrypted
- C. When the payload sent through the message broker must use XML format
- D. When the messaging broker must support point-to-point messaging

67. A company is designing a Mule application to consume batch data from a partner's FTPS server, `ftps.partner.com`. The data files have been compressed (but not encrypted) and then digitally signed by the partner using Pretty Good Privacy (PGP).

What inputs are required for the application to securely login to the FTPS server and to consume these files?

- A. The company's FTPS server login username and password

A TLS context trust store containing a public certificate for the company

The company's PGP public key that was used to sign the files

- B. The company's FTPS server login username and password

A TLS context trust store containing a public certificate for `ftps.partner.com`

The partner's PGP public key that was used to sign the files

- C. The partner's PGP public key used by the company to login to the FTPS server

A TLS context key store containing the private key for the company

The partner's PGP private key that was used to sign the files

- D. The partner's PGP public key used by the company to login to the FTPS server

A TLS context key store containing the private key for ftps.partner.com

The company's PGP private key that was used to sign the files

68. An application deployed in a CloudHub production environment is throwing intermittent exceptions. To troubleshoot the issue, a developer wants to enable DEBUG level logging in production for the org.apache.cxf package in the deployed Mule application.

How should the developer update the logging configuration in order to enable this package specific debugging?

- A. In the Mule application's log4j2.xml file, change the root logger's level property to DEBUG, then redeploy the Mule application to the CloudHub production environment

- B. In Anypoint Runtime Manager, in the Deployed Application Properties tab for the Mule application, add a line item with DEBUG level for package org.apache.cxf and apply the changes

- C. In the Mule application's log4j2.xml file, add an AsyncLogger element with name property set to org.apache.cxf and level set to DEBUG, then redeploy the Mule application in the CloudHub production environment

- D. In Anypoint Monitoring, define a logging search query with class property set to org.apache.cxf and level set to DEBUG

69. A manufacturing company wants to share inventory updates with dealers D1 and D2 asynchronously and concurrently via Queues Q1 and Q2. Dealer D1 must consume the messages from the Queue Q1, and Dealer D2 must consume the messages from the Queue Q2.

Dealer D1 has implemented a retry mechanism to reprocess the transaction in case of any errors while processing the inventory updates. Dealer D2 has not implemented any retry mechanism.

How should the dealers acknowledge the message to avoid message loss and minimize impact on the current implementation?

- A. Dealer D1 can use AUTO acknowledgement and Dealer D2 must use MANUAL acknowledgment and acknowledge the message after successful processing

- B. Dealer D1 can use AUTO acknowledgement and Dealer D2 can use IMMEDIATE acknowledgment and acknowledge the message after successful processing
- C. Dealer D1 and Dealer D2 must use AUTO acknowledgement and acknowledge the message after successful processing
- D. Dealer D2 must use AUTO acknowledgement and Dealer D1 can use MANUAL acknowledgment and acknowledge the message after successful processing

70. A corporation has deployed Mule applications to different customer-hosted Mule runtimes. Mule applications deployed to these Mule runtimes are managed by Anypoint Platform.

What needs to be installed or configured (if anything) to monitor these Mule applications from Anypoint Monitoring, and how is monitoring data from each Mule application sent to Anypoint Monitoring?

- A. Enable monitoring of individual Mule applications from the Runtime Manager application settings.

Runtime Manager sends monitoring data to Anypoint Monitoring for each deployed Mule application.

- B. Install an Anypoint Monitoring agent on each Mule runtime.

Each Anypoint Monitoring agent sends monitoring data from the Mule applications running in its Mule runtime to Anypoint Monitoring.

- C. Leave the out-of-the-box Anypoint Monitoring agent unchanged in its default Mule runtime installation.

Each Anypoint Monitoring agent sends monitoring data from the Mule applications running in its Mule runtime to Runtime Manager, then Runtime Manager sends monitoring data to Anypoint Monitoring.

- D. Install a Runtime Manager agent on each Mule runtime.

Each Runtime Manager agent sends monitoring data from the Mule applications running in its Mule runtime to Runtime Manager, then Runtime Manager sends monitoring data to Anypoint Monitoring.

71. When a Mule application using VM queues is deployed to a customer-hosted cluster or multiple CloudHub v1.0 workers/replicas, how are messages consumed across the nodes?

- A. Sequentially, only from the primary node
- B. In a non-deterministic way

- C. Round-robin, within an XA transaction
- D. Sequentially, from a dedicated Anypoint MQ queue

72. An organization plans to use the Anypoint Platform audit logging service to log Anypoint MQ actions.

What consideration must be kept in mind when leveraging Anypoint MQ Audit Logs?

- A. Anypoint MQ Audit Logs include logs for failed Anypoint MQ operations
- B. Anypoint MQ Audit Logs include logs for queue create, delete, modify, and purge operations
- C. Anypoint MQ Audit Logs include logs for sending, receiving, or browsing messages

73. A healthcare provider is implementing Anypoint Platform to manage its application infrastructure and is using the customer-hosted runtime for its business due to certain financial requirements it must meet. It has built a number of synchronous APIs and is currently hosting these on a Mule runtime on one server. These applications make use of a number of components, including heavy use of Object Stores and VM Queues.

Business has grown rapidly in the last year, and the healthcare provider is starting to receive reports of reliability issues from its applications. The DevOps team indicates that the APIs are currently handling too many requests, and this is overloading the server. The team has also mentioned that there is significant downtime when the server is down for maintenance.

What can the team add or upgrade to mitigate these issues?

- A. Add additional servers in a server group configuration
- B. Upgrade the networking hardware of the server
- C. Upgrade the CPU and memory of the server
- D. Add a load balancer and add additional servers in a cluster configuration

74. A bank is implementing a REST API in a Mule application to receive an array of accounts from an online banking platform user interface (UI), retrieve account balances for those accounts from a backend Finance system, and then return the account balances so they can be displayed in the online banking platform UI. As part of the processing, the MuleSoft API also needs to insert the retrieved account data into an Audit Database for auditing purposes, the auditing process should not add latency to the account balance retrieval response back to the online banking platform UI.

The retrieveBalances flow in the Mule application is designed to use an operation in a connector to the Finance system (the Finance operation) that can only look up one account record at a time, and a operation from a different connector to the Audit system (the Audit operation) that can only insert one account record at a time.

To best meet the performance-related requirements, what scope or scopes should be used and how should they be used to incorporate the Finance operation and Audit operation into the retrieveBalances flow?

- A. Wrap the Finance operation in a Until-Successful scope.

Wrap the Audit operation in a Try-Catch scope.

- B. Wrap both connector operations in a For-Each scope.

- C. Wrap the Finance operation in a Parallel For-Each scope.

Wrap the Audit operation in a Async scope.

- D. Wrap both connector operations in a Async scope.

75. An organization's security requirements mandate centralized control at all times over authentication and authorization of external applications when invoking web APIs managed on Anypoint Platform.

What Anypoint Platform feature is the most straightforward and maintainable to use to meet this requirement?

- A. Identify management configured in Access Management

- B. Enterprise Security module coded in Mule applications

- C. Client management configured in Access Management

- D. External access configured in API Manager

76. A Mule application contains a Batch Job scope with several Batch Step scopes. The Batch Job scope is configured with a batch block size of 25. A payload with 4,000 records is received by the Batch Job scope.

When there are no errors, how does the Batch Job scope process records within and between the Batch Step scopes?

- A. The Batch Job scope processes multiple record blocks in parallel, and a block of 25 records can jump ahead to the next Batch Step scope over an earlier block of records.

Each Batch Step scope is invoked with one record in the payload of the received Mule event.

For each Batch Step scope, all 25 records within a block are processed in parallel.

All the records in a block must be completed before the block of 25 records is available to the next Batch Step scope.

- B. The Batch Job scope processes multiple record blocks in parallel, and a block of 25 records can jump ahead to the next Batch Step scope over an earlier block of records.

Each Batch Step scope is invoked with one record in the payload of the received Mule event.

For each Batch Step scope, all 25 records within a block are processed sequentially.

All the records in a block must be completed before the block of 25 records is available to the next Batch Step scope.

- C. The Batch Job scope processes multiple record blocks in parallel.

Each Batch Step scope is invoked with a batch of 25 records in the payload of the received Mule event.

For each Batch Step scope, all 4,000 records are processed in parallel.

Individual records can jump ahead to the next Batch Step scope before the rest of the records finish processing

77. An organization plans to extend its Mule APIs to the EU (Frankfurt) region. Currently, all Mule applications are deployed to CloudHub 1.0 in the default North American region, from the North America control plane, following this naming convention: `**(API-name)**-**(environment)**` (for example, `Orders-sapi -dev`, `Orders-sapi -qa`, `Orders-sapi -prod`, etc.).

There is no network restriction to block communications between APIs.

What strategy should be implemented in order to deploy the same Mule APIs to the CloudHub 1.0 EU region from the North America control plane, as well as to minimize latency between APIs and target users and systems in Europe?

- A. In API Manager, leave the Region property blank (default) to deploy an API proxy named `**(API-name)**-**proxy**-**(environment).de-c1` for each Mule application.

- Communicate the new url `**(API-name)**-**proxy**-**(environment).de-c1.cloudhub.io` to the consuming API clients in Europe.

- B. In API Manager, set the Region property to EU (Frankfurt) to create an API proxy named ****(API-name)**-**proxy**-(environment)**** for each Mule application.

- Communicate the new url ****(API-name)**-**proxy**-(environment).de-c1.cloudhub.io** to the consuming API clients in Europe.

- C. In Runtime Manager, for each Mule application deployment, set the Region property to EU (Frankfurt) and reuse the same Mule application name as in the North American region.

- Communicate the new urls ****(API-name)**-(environment).de-c1.cloudhub.io** to the consuming API clients in Europe.

- D. In Runtime Manager, for each Mule application deployment, leave the Region property blank (default) and change the Mule application name to ****(API-name)**-(environment).de-c1**.

- Communicate the new urls ****(API-name)**-(environment).de-c1.cloudhub.io** to the consuming API clients in Europe.

78. A Mule application is designed to call the Google Maps API to perform a distance computation, and is deployed to CloudHub.

At a minimum, what (if anything) should be configured in the TLS context of the HTTP Request configuration to meet these requirements?

- A. The configuration is built-in and nothing extra is required for the TLS context

- B. Download the Google public certificate from a browser, generate a JKS file from it, and add it in TrustStore as part of the TLS context

- C. Download the Google public certificate from a browser, generate a JKS file from it, and add it in KeyStore as part of the TLS context

- D. Request a private key from Google, create a PKCS#12 file with it, and add it in KeyStore as part of the TLS context

79. A leading eCommerce giant will use MuleSoft APIs on Runtime Fabric (RTF) to process customer orders. Some customer-sensitive information, such as credit card information, is required in request payloads or is included in response payloads in some of the APIs. Other API requests and responses are not authorized to access some of this customer-sensitive information but have been implemented to validate and transform based on the structure and format of this customer-sensitive information (such as account IDs, phone numbers, and postal codes).

What approach configures an API gateway to hide sensitive data exchanged between API consumers and API implementations, but can convert tokenized fields back to their original value for other API requests or responses, without having to recode the API implementations?

- A. Create a masking format and use it to apply a tokenization policy in an API gateway to mask sensitive values in message payloads with "****" characters, and apply a corresponding detokenization policy to return the original values to other APIs
- B. Use a field-level encryption policy in an API gateway to replace sensitive fields in message payload with encrypted values, and apply a corresponding field-level decryption policy to return the original values to other APIs
- C. Create both masking and tokenization formats and use both to apply a tokenization policy in an API gateway to mask sensitive values in message payloads with "****" characters, and apply a corresponding detokenization policy to return the original values to other APIs
- D. Create a tokenization format and use it to apply a tokenization policy in an API gateway to replace sensitive...

80. An architect is designing a Mule application to meet the following two requirements:

1. The application must process files asynchronously and reliably from an FTPS server to a back-end database using VM intermediary queues for load-balancing Mule events.
2. The application must process a medium rate of records from a source to a target system using a Batch Job scope.

To make the Mule application more reliable, the Mule application will be deployed to two CloudHub 1.0 workers.

Following MuleSoft-recommended best practices, how should the Mule application deployment typically be configured in Runtime Manager to best support the performance and reliability goals of both the Batch Job scope and the file processing VM queues?

- A. Check the Persistent VM queues checkbox in the application deployment configuration
- B. In the Runtime Manager Properties tab, disable persistent VM queues for Batch Job scopes
- C. In the Runtime Manager Properties tab, enable persistent VM queues for the FTPS connector
- D. Check the Non-persistent VM queues checkbox in the application deployment configuration

81. An organization has strict unit test requirements that mandate every Mule application must have an MUnit test suite with a test case defined for each flow and a minimum test coverage of 80 percent.

A developer is building an MUnit test suite for a newly developed Mule application that sends API requests to an external REST API.

What is an effective approach for successfully executing the MUnit tests of this Mule application while still achieving the required test coverage for the MUnit tests?

- A. Mock the REST API invocations in the MUnit tests, and then forward mocked requests to the external REST API

- B. Mock the REST API invocations in the MUnit tests, and then return a mock response for those invocations

- C. Create a mocking service flow to simulate standard responses from the REST API, and then configure the Mule flows to call the mocking service flow

- D. Invoke the external endpoint of the REST API from the Mule flows

82. A project team is working on an API implementation using an API RAML specification as a starting point. The team has updated the RAML specification to include new operations, and has published a new version to Exchange. Meanwhile, another team is working on an API client Mule application to consume the same API implementation.

During this development, what must be performed by the Mule application team to take advantage of the newly added operations?

- A. Scaffold the API client application with APIkit using the new RAML specification

- B. Update the REST connector from Exchange in the API client application

- C. Update the API connector in the API implementation and publish to Exchange

- D. Copy the API implementation's updated APIkit component into the API client

83. One of the back-end systems invoked by the API implementation enforces rate limits on the number of requests a particular client can make. Both the back-end system and the API implementation are deployed to several non-production environments, including the staging environment, and to a production environment.

Rate limiting of the back-end system applies to all non-production environments. The production environment, however, does not have any rate limiting.

What is the most effective approach to conduct performance tests of the API implementation in the non-production staging environment?

Staging environment

(Performance test scenario)

Performance testing client

Mule API implementation

Rate limiting applied

Back-end system

- A. Conduct scaled-down performance tests in the staging environment against the rate-limited back-end system. Then, upscale performance results to full production scale.
- B. Include logic within the API implementation that bypasses invocations of the back-end system in the staging environment, and invoke a mocking service that replicates typical back-end system responses. Then, conduct performance tests using this API implementation.
- C. Create a mocking service that replicates the back-end system's production performance characteristics. Then, configure the API implementation to use the mocking service and conduct the performance tests.
- D. Use MUnit to simulate standard responses from the back-end system. Then, conduct performance tests to identify other bottlenecks in the system.

84. A cluster with two customer-hosted Mule runtimes is hosting an application that has a flow with a JMS listener configured to subscribe to messages from a topic destination.

How should the JMS listener's `primaryNodeOnly` parameter and subscription parameter be configured so the JMS listener can receive messages in all the nodes of the cluster but only process each message once?

- A. Use the parameter `primaryNodeOnly="false"` on the JMS listener with a shared subscription
- B. Use the parameter `primaryNodeOnly="true"` on the JMS listener with a non-shared subscription
- C. Use the parameter `primaryNodeOnly="true"` on the JMS listener with a shared subscription
- D. Use the parameter `primaryNodeOnly="false"` on the JMS listener with a non-shared subscription

85. A gaming company has implemented an API as a Mule application and deployed the API implementation to a CloudHub 2.0 private space. The API implementation must connect to a

mainframe application running in the customer's on-premises corporate data center and also to a Kafka cluster running in an Amazon AWS VPC.

What is the most efficient way to enable the API to securely connect from its private space to the mainframe application and Kafka cluster?

- A. In Runtime Manager, configure an Anypoint VPN to route from the CloudHub 2.0 private space to the on-premises data center. In the AWS account, attach the private space to an AWS transit gateway that routes to the AWS VPC.

- B. In Runtime Manager, set up VPC peering between the CloudHub 2.0 private network and the on-premises data center.

In the AWS account, set up VPC peering between the AWS VPC and the CloudHub 2.0 private network.

- C. In the AWS account, attach the private space directly to the AWS VPC.

In the AWS account, use an AWS transit gateway to route from the AWS VPC to the on-premises data center.

- D. In the AWS account, attach the CloudHub 2.0 private space to an AWS transit gateway that routes from the CloudHub 2.0 private space to the on-premises data center. In Runtime Manager, configure an Anypoint VPN to route from the CloudHub 2.0 private space to the AWS VPC.

86. A manufacturing company plans to deploy Mule applications to its own Azure Kubernetes service infrastructure. The organization wants to make the Mule applications more available and robust by deploying each Mule application to an isolated Mule runtime in a Docker container while managing all the Mule applications from the MuleSoft-hosted control plane.

What choice of runtime plane meets these organizational requirements?

- A. CloudHub 2.0

- B. Anypoint Runtime Fabric

- C. Anypoint Service Mesh

- D. Customer-hosted self-provisioned runtime plane

87. An organization is successfully using API-led connectivity. However, as the application network grows, all the manually performed tasks to publish, share and discover, register, apply policies to, and

deploy an API are becoming repetitive, which is driving the organization to automate this process using efficient CI/CD pipelines.

Considering Anypoint Platform's capabilities, how should the organization automate its API lifecycle?

- A. Use Anypoint CLI or Anypoint Platform REST APIs with a scripting language such as Groovy
- B. Use Runtime Manager REST APIs for API management and Maven for API deployment
- C. Use Maven with a custom configuration required for the API lifecycle
- D. Use Exchange REST APIs for API management and Maven for API deployment

88. An organization is using MuleSoft to develop APIs and automatically deploy them to both CloudHub and on-premises targets. Recently, it has decided to include Runtime Fabric deployment targets as well, and Mule infrastructure is set up for this option.

What can now be used to automate deploying Mule applications to all three types of deployment targets?

- A. Anypoint Monitoring agent configuration file
- B. The Object Store v2 REST API
- C. The Runtime Manager agent configuration file
- D. The Mule Maven plugin

89. Refer to the exhibits.

An Order microservice and a Fulfillment microservice are being designed to communicate with their clients through message-based integration (and NOT through API invocations).

The Order microservice publishes an Order message (a kind of command message) containing the details of an order to be filled. The intent is that Order messages are only consumed by one Mule application: the Fulfillment microservice.

The Fulfillment microservice consumes Order messages, fills the orders described therein, and then publishes a Fulfillment message (a kind of event message). Each Fulfillment message can be consumed by any interested Mule application, and the Order microservice is one such Mule application.

What is the recommended choice of message broker(s) and message destination(s) in this scenario?

- A. Order messages are sent to an Anypoint MQ exchange.
- Fulfillment messages are sent to an Anypoint MQ queue.

Both microservices interact with Anypoint MQ as the message broker, which must therefore scale to support the

- B. Order messages are sent to a JMS topic.
- Fulfillment messages are sent to a JMS queue.

The Order microservice interacts with one JMS provider (message broker) and the Fulfillment microservice interacts with a different JMS provider so that both message brokers can be chosen and scaled to best support the load of each microservice.

- C. Order messages are sent to a JMS queue.
- Fulfillment messages are sent to a JMS topic.

Both microservices interact with the same JMS provider (message broker) instance, which must therefore scale to support the load of both microservices.

90. An organization is in the process of building automated deployments using a CI/CD process. As a part of automated deployments, it wants to apply policies to API instances.

What tool can the organization use to promote and deploy API Manager policies?

- A. Runtime Manager agent
- B. Mule Maven plugin
- C. Anypoint CLI
- D. MUnit Maven plugin

91. A healthcare client must connect its Mule APIs that are deployed to CloudHub 1.0 to systems in an existing on-premises datacenter through private and secure connections.

Which network architecture supports these security requirements to connect the Mule APIs to the on-premises datacenter?

- A. Static IP addresses for the Mule applications deployed to the CloudHub 1.0 Shared Worker Cloud
- Matching IP included listing in the load balancer and firewall rules in the VPC and on-premises network

- B. Static IP addresses for the Mule applications deployed to the CloudHub 1.0 Shared Worker Cloud
Matching firewall rules and IP includedlisting in the on-premises network

- C. An Amypoint VPC connected to the on-premises network using a dedicated load balancer
Matching IP includedlisting in the load balancer and firewall rules in the VPC and on-premises network

- D. An Amypoint VPC connected to the on-premises network using an iPsec tunnel
Matching firewall rules in the VPC and on-premises network

92. Why would an Enterprise Architect use a single enterprise-wide canonical data model (CDM) when designing an integration solution using Anypoint Platform?

- A. To remove the need to perform data transformation when processing message payloads in Mule applications
- B. To automate AI-enabled API implementation generation based on normalized backend databases from separate vendors
- C. To leverage a data abstraction layer that shields existing Mule applications from non-backward compatible changes to the model's data structure

- D. To reduce dependencies when integrating multiple systems that use different data formats

93. A payment processing company has implemented a Payment Processing API Mule application to process credit card and debit card transactions. Because the Payment Processing API handles highly sensitive information, the payment processing company requires that data must be encrypted both in-transit and at-rest.

To meet these security requirements, consumers of the Payment Processing API must create request message payloads in a JSON format specified by the API, and the message payload values must be encrypted.

How can the Payment Processing API validate requests received from API consumers?

- A. A Transport Layer Security (TLS) - Inbound policy can be applied in API Manager to decrypt the message payload and the Mule application implementation can then use the JSON Validation module to validate the JSON data

- B. The Mule application implementation can use the APIkit module to decrypt and then validate the JSON data

- C. The Mule application implementation can use the Validation module to decrypt and then validate the JSON data
- D. The Mule application implementation can use DataWeave to decrypt the message payload and then use the JSON Scheme Validation module to validate the JSON data

94. An organization has deployed Runtime Fabric on a single server configured with an eight-node cluster using the performance profile. An API uses a non-persistent Object Store for maintaining state data from deployed Mule applications.

What will be the impact to the state data if the server crashes?

- A. State data is preserved
- B. State data is lost
- C. State data is rolled back to a previously saved version
- D. State data is duplicated

95. An organization is building out a test suite for their application using MUnit. The Integration Architect has recommended using Test Recorder in Anypoint Studio to record the processing flows and then configure unit tests based on the captured events.

What is a core consideration that must be kept in mind while using Test Recorder?

- A. Tests for flows cannot be created if Mule errors are raised inside the flows, even if the errors are handled by On-Error Continue error handlers
- B. The Recorder supports mocking a message before or inside a foreach processor
- C. The Recorder supports loops where the structure of the data being tested changes inside the iteration
- D. Mocking values resulting from parallel processes are possible and will not affect the execution of the processors that follow in the test

96. An organization has decided on a cloud migration strategy to minimize the organization's own IT resources. Currently, the organization has all of its Mule applications running on its own premises and uses an on-premises load balancer that exposes all APIs under the base URL <https://api.acme.com>.

As part of the migration strategy, the organization plans to migrate all of its Mule applications and the on-premises load balancer to CloudHub 1.0.

What is the most straightforward and cost-effective approach to Mule application deployment and load balancing that preserves the public URLs?

- A. Deploy the Mule applications to CloudHub 1.0.

Update the CNAME record for api.acme.com in the organization's DNS server to point to the A record of the CloudHub Shared Load Balancer (SLB).

Apply mapping rules in the SLB to map URLs to their corresponding Mule applications.

- B. Deploy the Mule applications to CloudHub 1.0.

Update the CNAME record for api.acme.com in the organization's DNS server to point to the A record of a CloudHub 1.0 Dedicated Load Balancer (DLB).

Apply mapping rules in the DLB to map URLs to their corresponding Mule applications.

C. For each migrated Mule application, deploy an API proxy Mule application to CloudHub 1.0, with all traffic to the Mule applications routed through a CloudHub 1.0 Dedicated Load Balancer (DLB).

Update the CNAME record for api.acme.com in the organization's DNS server to point to the A record of the CloudHub 1.0 DLB.

Apply mapping rules in the DLB to map each API proxy application to its corresponding Mule application.

D. Deploy the Mule applications to CloudHub 1.0.

97. An organization has built a large monolithic application over the years and is currently looking to transition to a microservices architecture. During this transition phase, it has built some System APIs as Mule applications that provide access to different pieces of functionality of the monolithic application. The System APIs are deployed to Mule runtimes in a customer-hosted EC2 instance in AWS.

The System APIs communicate directly with one database used by the monolithic application. Currently, a Database connector and configuration for that connector must be built into every project. There is a plan to partition this database into various smaller data stores in the future.

What is the MuleSoft-recommended best practice to share the connector and configuration information among the APIs?

- A. Build another System API that connects to the database, and refactor all the other APIs to make requests through the new System API to access the database

- B. Create an API proxy for each System API and share the Database connector configuration with all the API proxies via an automated policy

- C. Build a separate Mule domain project for each API, and configure each of them to use a file on a shared file store to load the configuration information dynamically

- D. Build a Mule domain project, add the Database connector and configuration to it, and reference this one domain project from each System API

98. Which Salesforce API is invoked to deploy, retrieve, create, update, or delete customization information, such as custom object definitions using Mule Salesforce Connectors in a Mule application?

- A. sObject Platform Action API
- B. Process Rules API
- C. User Interface API
- D. Metadata API

99. Which operation can be performed through a JMX agent enabled in a Mule application?

- A. Set a particular log#j2 log level to TRACE
- B. View object store entries
- C. Deploy a Mule application
- D. Replay an unsuccessful message

100. Refer to the exhibit.

A manufacturing company is developing a new set of APIs for its retail business. One of the APIs is a Master Look Up API, which is a System API. The API uses a persistent object-store. This API will be used by almost all other APIs to provide master lookup data.

The Master Look Up API is deployed on two CloudHub workers of 0.1 vCore each because there is a lot of master data to be cached. Most of the master lookup data is stored as a key-value pair. The cache gets refreshed if the key is not found in the cache.

During performance testing, it was determined that the Master Look Up API has a high response time due to the latency of database queries executed to fetch the master lookup data.

What two methods can be used to resolve these performance issues?

- A. Implement locking to synchronize access to the Object Store
- B. Upgrade the vCore size from 0.1 vCore to 0.2 vCore
- C. Implement the HTTP caching policy for all GET endpoints for the Master Look Up API

101. A new Mule application has been deployed through Runtime Manager to CloudHub 1.0 using a CI/CD pipeline with sensitive properties set as cleartext into Runtime Manager Administrator opened a high priority incident ticket about this violation of their security requirements indicating these sensitive properties values must not be stored or visible in Runtime Manager but should be changeable in Runtime Manager by Administrators with proper permissions.

How can the Mule application be deployed while safely hiding the sensitive properties?

- A. Create a variable for each sensitive property and declare them as hidden in the CI/CD pipeline scripts
- B. Add a new `wrapper.java.additional.xx` parameter for each sensitive property in the `wrapper.conf` file used by the CI/CD pipeline scripts
- C. Add encrypted versions of the sensitive properties as global configuration properties in the Mule application
- D. Add an ArrayList of all the sensitive properties' names in the `mule-artifact.json` file of the application

102. A Mule application uses an HTTP Request operation to invoke an external API. The external API follows the HTTP specification for proper status code usage.

What is a possible cause when a 3xx status code is returned to the HTTP Request operation from the external API?

- A. The request was received and processed by the external API
- B. The request was not received by the external API
- C. The request was redirected to a different URL by the external API
- D. The request was received and rejected by the external API

103. A Mule application is required to periodically process a large data set from a back-end database to Salesforce CRM using a Batch Job scope configured properly to process the high rate of records. The Mule application is deployed to two CloudHub 1.0 workers with the persistent VM queues option disabled (unchecked).

What is the consequence if a worker that is processing records from a Batch Job scope crashes during records processing?

- A. About half the remaining records will be processed by the second worker and the rest will be lost, resulting in some unprocessed records

- B. Remaining records will be lost, resulting in unprocessed records

- C. All the remaining records will be processed by a new replacement worker, resulting in “exactly once” processing of every record

- D. The entire batch job will be processed from scratch by the second worker, resulting in some duplicate processing

104. A marketing organization is designing a Mule application to process campaign data. The Mule application will periodically check for a file in an SFTP location and process the records in the file. The size of the file can vary from 10MB to 50B. Due to the limited availability on/<core>, the Mule application is deployed to a single CloudHub worker/replica configured with vCore size 0.2.

The application must transform and send different formats of this file to three different downstream SFTP locations.

What is the most direct way to configure the SFTP operations or event sources to process the large files to support these deployment requirements?

- A. Use a file-stored repeatable stream

- B. Use an in-memory nonrepeatable stream

- C. Use a file-stored nonrepeatable stream

- D. Use an in-memory repeatable stream

105. An airline is architecting an API-led connectivity project to integrate its flight data into an online aggregation website. The connections must allow for secured communication, high performance, and asynchronous message exchange.

What are suitable interface technologies for the connectors in this integration assuming that Mulesoft fully supports those technologies and that Anypoint Connectors exist for these interfaces?

- A. AsyncAPI over SGS AMOIID with RabbitMX JSON/REST over HTTPS

- B. XML over ActiveMQ XML over FTP XML/REST over HTTPS

- C. SCAP over HTTPS HDP over TLS gRPC over HTTPS

- D. CSV over FTP YAML over TLS JSON over HTTPS

106. In a Mule 4 application that will be deployed to customer-hosted Mule runtimes, the Java module's Invoice Static operation needs to invoke a static method in a dependency JAR file.

How can the dependency be made visible to the Java module's classloader?

- A. Use a Maven command to include the dependency JAR file when packaging the application
- B. In the Java Invoice Static module configuration, configure the path and name of the dependency JAR file
- C. Add the dependency JAR file to the Java classpath by setting JVM parameters
- D. Configure the dependency as a shared library in the project's POM file

107. Following MuleSoft's recommended best practices for API governance and API policies, a project team has used RAML specifications to document and publish functional requirements and detailed design definitions of its APIs. These API specifications have been used by various stakeholders to implement APIs.

Later, the project team requires all API specifications to be augmented with an additional non-functional requirement (NFR) to protect the backend services from a high rate of requests, according to defined service-level agreements (SLAs). The NFRs SLAs are based on a new tiered subscription level "Golet", "Silver", or "Platinum" that must be tied to a new parameter that is being added to the Accounts object in their enterprise data model.

Following MuleSoft's recommended best practices, how should the project team now convey the necessary non-functional requirement to stakeholders?

- A. Update each API specification with a shared RAML fragment required to implement the NFR and publish the RAML fragment and each modified API specification to Exchange
- B. Create a shared RAML fragment required to implement the NFR, list each API implementation endpoint in the RAML fragment, and publish the RAML fragment to Exchange
- C. Create and deploy API proxies in API Manager for the NFR, change the baseUrl in each API specification to the corresponding API proxy implementation endpoint, and publish each modified API specification to Exchange
- D. Update each API specification with comments about the NFRs SLAs and publish each modified API specification to Exchange

108. A Mule application is synchronizing customer data between two different database systems.

What is a main benefit of using eXtended Architecture (XA) transactions rather than local transactions to synchronize these two different database systems?

- A. An XA transaction automatically rolls back operations against both database systems if any operation fails

- B. An XA transaction synchronizes the database systems with the least amount of Mule configuration or coding

- C. An XA transaction handles the largest number of requests in the shortest time

- D. An XA transaction writes to both database systems quickly and simultaneously

109. An organization is designing multiple Mule applications to run on CloudHub in a single Anypoint VPC and that must share data using a common persistent Anypoint Object Store v2 (OSv2).

What design gives these Mule applications access to the same Object Store instance?

- A. The Object Store v2 can be shared across CloudHub applications with a configured OS v2 connector

- B. The Object Store v2 REST API configured to access the persistent Object Store

- C. A VM connector configured to directly access the persistence queue of the persistent Object Store

- D. An Anypoint MG connector configured to directly access the persistent Object Store

110. A company is designing an Integration Mute application to process orders by submitting them to a back-end system for offline processing. Each order will be received by the Mute application through an HTTP POST and must be acknowledged immediately. Once acknowledged, the order will be submitted to a back-end system. Orders that cannot be successfully submitted due to rejections from the back-end system will need to be processed manually (outside the back-end system).

The Mute application will be deployed to a customer-hosted runtime and will be able to use an existing ActiveMQ broker if needed. The ActiveMQ broker is located inside the organization's firewall. The back-end system has a track record of unreliability due to both minor network connectivity issues and longer outages.

What combination of Mute application components and ActiveMQ queues are required to ensure automatic submission of orders to the back-end system while supporting but minimizing manual order processing?

- A. An Until Successful scope to call the back-end system

One or more ActiveMQ long retry queues

One or more ActiveMQ dead letter queues for manual processing

- B. One or more On Error scopes to assist calling the back-end system

An Until Successful scope containing VM components for long retries

A persistent dead-letter VM queue configured in CloudHub

- C. One or more On Error scopes to assist calling the back-end system

One or more ActiveMQ long retry queues

A persistent dead-letter Object Store configured in the CloudHub Object Store service

- D. A Batch Job scope to call the back-end system

An Until Successful scope containing Object Store components for long retries

A dead-letter object store configured in the Mule application

111. A Mule solution implements an Experience API and System API as Mule applications that process messages in this order:

1. The Experience API receives a request Anypoint MQ message (REQU) with a payload containing a variable-length list of request objects.

2. The Experience API uses a For Each scope to split the list into individual objects and sends each object as a REQU_OBJECT message to an

Anypoint MQ queue named REQUEST.

3. The System API listens on the REQUEST queue, processes each message independently of all other messages, and sends a response message to a response queue named RESPONSE_OBJECT. The System API processing times can vary and are occasionally fairly long.

4. The Experience API listens on the RESPONSE_OBJECT queue and receives, transforms, and accumulates each response.

5. After receiving all the responses corresponding to the original REQU message, the Experience API creates and publishes a response Anypoint MQ message (RESP) to an Anypoint MQ queue named RESPONSE with a payload containing the list of responses returned by the System API, organized in the same order as the request objects originally sent in the REQU message.

Assume successful response messages are returned by Experience API to the RESPONSE queue for all REQU messages.

What is required so the Experience API can ensure the length and order of the list of objects in RESP and REQU match, while at the same time maximizing message throughput?

- A. Perform all communication involving the System API synchronously from within the For Each scope so that responses in RESPONSE_OBJECT are in the exact same order as request objects in the

RESP message. Within the For Each scope, use an Async scope configured with a persistent object store.

- B. Use a Scatter-Gather within the For Each scope to ensure response message order. Configure the Scatter-Gather with a persistent object store.

- C. After the For Each scope, add a second For Each scope configured with a persistent object store to collect response messages in the order in which they arrive, and then send the RESP message using this list of responses.

- D. Use a persistent object store to keep track of the list length and all object indexes in the REQU message, both in the For Each scope and in all communication involving the System API. Set the maxConcurrency parameter for the receiving flow in the Experience API to a value greater than 1.

112. A stock trading company handles millions of trades a day and requires excellent performance and reliability within its stock trading system. The company operates a number of event-driven APIs implemented as Mule applications that are hosted on various customer-hosted Mule clusters and needs to enable message exchanges between the APIs within their internal network using shared message queues.

What is an effective way to meet the cross-cluster messaging requirements of its event-driven APIs?

- A. Persistent VM queues with automatic acknowledgements

- B. extended Architecture (XA) transactions and XA connected components with manual acknowledgements

- C. JMS transactions with automatic acknowledgements

- D. Non-transactional JMS operations with a reliability pattern and manual acknowledgements

113. Refer to the exhibit.

In the flow of this Mule application, what thread pool handles requests to the HTTP Listener, and what thread pool handles the Logger component?

Thread pool Flow

Liftflow

Logger

Error handling

- A. BLOCKING_IO, UBER
- B. Shared NIO Selector Pool, CPU_LITE
- C. UBER, Dedicated NIO Selector Pool
- D. CPU_LITE, CPU_INTENSIVE

114. An organization plans to migrate all its Mule applications to Runtime Fabric (RTF). Currently, all Mule applications have been deployed to CloudHub using automated CI/CD scripts.

What steps should be taken to properly migrate the applications from CloudHub to RTF, while keeping the same automated CI/CD deployment strategy?

- A. A runtimeFabric dependency should be added as a mule-plugin to the pom.xml file in all the Mule applications.

A --runtimeFabric command-line parameter should be added to the CI/CD deployment scripts.

- B. A runtimeFabricDeployment profile should be added to Mule configuration properties
YAML files in all the Mule applications.

CI/CD scripts must be modified to use the new configuration properties.

- C. A runtimeFabricDeployment profile should be added to the pom.xml file in all the Mule applications.

CI/CD scripts must be modified to use the new RTF profile.

- D. The pom.xml and Mule configuration YAML files can remain unchanged in each Mule application.

A --runtimeFabric command-line parameter should be added to the CI/CD deployment scripts

115. A company is modernizing its legacy systems to accelerate access to applications and data while supporting the adoption of new technologies. The key to achieving this business goal is unlocking the company's systems and data, including a set of existing services hosted on-premises that can be accessed by authorized external clients. The IT staff is mainly experienced only with supporting its legacy systems.

Considering the current aggressive backlog and the project delivery requirements, the company wants to take a strategic approach in the first phase of its transformation projects by quickly

deploying APIs in Mule runtimes that are able to scale, connect to on-premise systems, and migrate as needed.

Following MuleSoft best practices, what MuleSoft runtime deployment option best meets the company's goals to begin its digital transformation journey?

- A. CloudHub runtimes
- B. Customer-hosted runtimes provisioned by a MuleSoft services partner
- C. Runtime Fabric on VMs/Base metal
- D. Customer-hosted self-provisioned runtimes

116. An organization plans to migrate its deployment environment from an on-premises cluster to a Runtime Fabric (RTF) cluster. The on-premises Mule applications are currently configured with persistent object stores.

There is a requirement to enable Mule applications deployed to the RTF cluster to store and share data across application replicas and through restarts of the entire RTF cluster.

How can these reliability requirements be met?

- A. Replace persistent object stores with persistent VM queues in each Mule application deployment
- B. Install the Object Store pod on one of the cluster nodes
- C. Configure the Persistence Gateway in the RTF installation
- D. Configure Anypoint Object Store v2 to share data between replicas in the RTF cluster

117. How does a timeout attribute in a configured transaction manager help inform design decisions while using a JMS connector listening for incoming messages in an eXtended Architecture (XA) transaction managed by the transaction manager?

- A. The timeout defines the time that is allowed to pass without the transaction ending explicitly, and after the timeout expires, the transaction rolls back
- B. After the timeout is exceeded, state JMS consumer threads are destroyed and new threads are created
- C. The timeout specifies the time allowed to pass between receiving JMS messages on the same JMS connection, and then after the timeout, a new JMS connection is established
- D. The time allowed to pass between committing the transaction and the completion of the Mule flow, and then after the timeout, flow processing triggers an error

118. A customer is running Mule applications on Runtime Fabric for Self-Managed Kubernetes (RTF-BYOK) in a multi-cloud environment.

Based on this configuration, how do Agents and Runtime Manager communicate, and what is exchanged between them?

- A. Runtime Manager Initiates traffic to Agents to check the status of clusters and auto-deploys new images when updates to Mult applications are available
- B. Runtime Manager communicates with Agents to provide updates on cluster status, download images, and update Mult applications (when available and chosen to do so by customer) from the container registry
- C. Agents communicate with Runtime Manager to provide updates on cluster status, download images, and update Mult applications (automatically when available) from the container registry
- D. Agents communicate with Runtime Manager to provide updates on cluster status, download images, and update Mult applications (when available and chosen to do so by customer) from the container registry

119. An organization is struggling with frequent plugin version upgrades and external plugin project dependencies. The team wants to minimize the impact on applications by creating best practices that will define a set of default dependencies across all new and in-progress projects.

How can these best practices be achieved with the applications having the least amount of responsibility?

- A. Create a Mule domain project with all the dependencies defined in its pom.xml file and add each application to the domain project
- B. Create a parent POM with all the required dependencies and reference it in each applications pom.xml file
- C. Add all the dependencies in each applications pom.xml file
- D. Create a Mule plugin project with all the dependencies and add it as a dependency in each applications pom.xml file

120. A cluster with two customer-hosted Mule runtimes is hosting an application that has a flow with a JMS listener configured to subscribe to messages from a topic destination.

How should the JMS listener's primaryNodeOnly parameter and subscription parameter be configured so the JMS listener can receive messages in all the nodes of the cluster but only process each message once?

- A. Use the parameter primaryNodeOnly="false" on the JMS listener with a shared subscription
- B. Use the parameter primaryNodeOnly="true" on the JMS listener with a non-shared subscription

- C. Use the parameter `primaryNodeOnly="true"` on the JMS listener with a shared subscription
- D. Use the parameter `primaryNodeOnly="false"` on the JMS listener with a non-shared subscription

Ravi-Hayatullah Team