

Question 1 of 60



Choose 1 option.

When can CloudHub Object Store v2 be used?

- ☒ A. To store an unlimited number of key-value pairs
- ☐ B. To store key-value pairs with keys up to 300 characters
- ☐ C. To store payloads with an average size greater than 15MB
- ☐ D. To store information in Mule 4 Object Store v1

Question 2 of 60



Choose 1 option.

A company deployed an API to a single worker/replica in the shared cloud in the U.S. West Region.

What happens when the Availability Zone experiences an outage?

- ☒ A. CloudHub will auto-redeploy the API in another Availability Zone in the U.S. West Region
- ☐ B. The Anypoint Platform admin is alerted when the API is experiencing an outage and needs to trigger the CI/CD pipeline to redeploy to the U.S. East Region
- ☐ C. The API will be unavailable until the availability comes back online, at which time the worker/replica will be auto-restarted
- ☐ D. CloudHub will auto-redeploy the API in the U.S. East Region

Question 3 of 60



Choose 1 option.

To minimize operation costs, a customer wants to use a CloudHub 1.0 solution. The customer's requirements are:

- Separate resources with two Business groups
- High-availability (HA) for all APIs
- Route traffic via Dedicated load balancer (DLBs)
- Separate environments into production and non-production

Which solution meets the customer's needs?

- ☐ A. One production and one non-production Virtual Private Cloud (VPC) per Business group.
Minimize CIDR aligning with projected application total.
Divide availability zones during deployment of APIs for HA.
- ☐ B. One production and one non-production Virtual Private Cloud (VPC).
Configure subnet to differentiate between business groups.
Allocate maximum CIDR per VPCs to make it easier to add Child groups.
Span VPC to cover three availability zones.
- ☐ C. One production and one non-production Virtual Private Cloud (VPC).
Use availability zones to differentiate between Business groups.
Allocate maximum CIDR per VPCs to ensure HA across availability zones
- ☒ D. One production and one non-production Virtual Private Cloud (VPC) per Business group.

Question 3 of 60



Choose 1 option.

- Route traffic via Dedicated load balancer (DLBs)
- Separate environments into production and non-production

Which solution meets the customer's needs?

- ☐ A. One production and one non-production Virtual Private Cloud (VPC) per Business group.
Minimize CIDR aligning with projected application total.
Divide availability zones during deployment of APIs for HA.
- ☐ B. One production and one non-production Virtual Private Cloud (VPC).
Configure subnet to differentiate between business groups.
Allocate maximum CIDR per VPCs to make it easier to add Child groups.
Span VPC to cover three availability zones.
- ☐ C. One production and one non-production Virtual Private Cloud (VPC).
Use availability zones to differentiate between Business groups.
Allocate maximum CIDR per VPCs to ensure HA across availability zones
- ☒ D. One production and one non-production Virtual Private Cloud (VPC) per Business group.
Minimize CIDR aligning with projected application total.
Choose a MuleSoft CloudHub 1.0 region with multiple availability zones.
Deploy multiple workers for HA.

Question 4 of 60



Choose 1 option.

An established communications company is beginning its API-led connectivity journey. The company has been using a successful Enterprise Data Model for many years. The company has identified a self-service account management app as the first effort for API-led, and it has identified the following APIs.

Experience layer: Mobile Account Management EAPI, Browser Account Management EAPI

Process layer: Customer Lookup PAPI, Service Lookup PAPI, Account Lookup PAPI

System layer: Customer SAPI, Account SAPI, Product SAPI, Service SAPI

According to MuleSoft's API-led connectivity approach, which API would not be served by the Enterprise Data Model?

- ☒ A. Mobile Account Management EAPI
- ☐ B. Service SAPI
- ☐ C. Customer SAPI
- ☐ D. Customer Lookup PAPI

Question 5 of 60



Choose 1 option.

Which statement is true about where an API policy is defined in Anypoint Platform and how it is applied to API instances?

- ☐ A. The API policy is defined in API Manager and then automatically applied to all API instances
- ☐ B. The API policy is defined in Runtime Manager as part of the API deployment to a Mule runtime and then only applied to the specific API instance
- ☒ C. The API policy is defined in API Manager for a specific API instance and then only applied to the specific API instance
- ☐ D. The API policy is defined in API Manager and then applied to all API instances at deployment time

Question 6 of 60



Choose 1 option.

A circuit breaker strategy is planned in order to meet the goal of improved response time and demand on a downstream API.

- Circuit Open: More than 10 errors per minute for three minutes
- Circuit Half-Open: One error per minute
- Circuit Closed: Less than one error per minute for five minutes

Out of several proposals from the engineering team, which option will meet this goal?

- ☐ A. Implement the strategy in a Mule application, and provide the settings in the YAML configuration
- ☐ B. Add the Circuit Breaker policy to the API instance, and configure the required settings
- ☐ C. Create Anypoint Monitoring alerts for Circuit Open/Closed configurations, and then implement a retry strategy for Circuit Half-Open configuration
- ☒ D. Create a custom policy that implements the circuit breaker and includes policy template expressions for the required settings

Question 7 of 60



Choose 1 option.

A developer for a transportation organization is implementing exactly one processing functionality in a Reservation Mule application to process and store passenger records. This Reservation application will be deployed to multiple CloudHub workers/replicas. It is possible that several external systems could send duplicate passenger records to the Reservation application.

An appropriate storage mechanism must be selected to help the Reservation application process each passenger record exactly once as much as possible. The selected storage mechanism must be shared by all the CloudHub workers/replicas in order to synchronize the state information to assist attempting exactly once processing of each passenger record by the deployed Reservation Mule application.

Which type of simple storage mechanism in Anypoint Platform allows the Reservation Mule application to update and share data between the CloudHub workers/replicas exactly once, with minimal development effort?

- ☐ A. Runtime Fabric Object Store
- ☐ B. In-memory Mule Object Store
- ☐ C. Non-persistent Object Store
- ☒ D. Persistent Object Store

Question 8 of 60



Choose 1 option.

An API implementation is deployed to CloudHub.

What conditions can be alerted on using the default Anypoint Platform functionality, where the alert conditions depend on the API invocations to an API implementation?

- ☐ A. When the API invocations are not over a secure TLS/SSL communication channel
- ☐ B. When the API invocations originate from a geography different than the API
- ☒ C. When the number of API invocations are below a threshold
- ☐ D. When the API invocations are sent directly to the internal DNS record of the API implementation

Question 9 of 60



Choose 1 option.

A developer from the Central IT team has created an initial version of the RAML definition in Design Center for an OAuth 2.0-protected System API and published it to Exchange. Another developer from LoB IT discovered the System API in Exchange and would like to leverage it in the Process API.

What is the MuleSoft-recommended approach for Process API to invoke the System API?

- ☒ A. The Process API uses the REST Connect Connector autogenerated in Exchange for the System API
- ☐ B. The Process API uses property YAML files to store the System API URLs and uses the HTTP Request Connector to invoke the System API
- ☐ C. The Process API needs to import an OAuth 2.0 module from Exchange first and update it with OAuth 2.0 credentials before the System API can be invoked
- ☐ D. The Process API manually updates the Process API POM file to include the System API as a dependency

Question 10 of 60



Choose 1 option.

A customer has an ELA contract with MuleSoft. An API deployed to CloudHub is consistently experiencing performance issues. Based on the root cause analysis, it is determined that autoscaling needs to be applied.

How can this be achieved?

- ☐ A. Configure a policy so that when the response time reaches a certain threshold the worker/replica type increases (vertical scaling)
- ☐ B. Configure two separate policies: When CPU and memory reach a certain threshold, increase the worker/replica type (vertical scaling) and the number of workers/replicas (horizontal scaling)
- ☒ C. Configure a policy based on CPU usage so that CloudHub auto-adjusts the number of workers/replicas (horizontal scaling)
- ☐ D. Configure a policy so that when the number of HTTP requests reaches a certain threshold the number of workers/replicas increases (horizontal scaling)

Question 11 of 60



Choose 1 option.

An application updates an inventory running only one process at any given time to keep the inventory consistent. This process takes 200 milliseconds (.2 seconds) to execute; therefore, the scalability threshold of the application is five requests per second.

What is the impact on the application if horizontal scaling is applied, thereby increasing the number of Mule workers?

- ☐ A. The application scalability threshold is now 10 requests per second
- ☒ B. The application scalability threshold is five requests per second regardless of the horizontal scaling
- ☐ C. The total process execution time is now 100 milliseconds (.1 seconds)
- ☐ D. Horizontal scaling cannot be applied to an already-running application

Question 12 of 60



Choose 1 option.

A Mule application implements an API. The Mule application has an HTTP Listener whose connector configuration sets the HTTPS protocol and hard-codes the port value. The Mule application is deployed to an Anypoint VPC and uses the CloudHub 1.0 Shared Load Balancer (SLB) for all incoming traffic.

Which port number must be assigned to the HTTP Listener's connector configuration so that the Mule application properly receives HTTPS API invocations routed through the SLB?

- ☒ A. 8082
- ☐ B. 80
- ☐ C. 8092
- ☐ D. 443

Question 13 of 60



Choose 1 option.

An enterprise is embarking on the API-led digital transformation journey, and the central IT team has started to define System APIs. Currently there is no Enterprise Data Model being defined within the enterprise, and the definition of a clean Bounded Context Data Model requires too much effort.

According to MuleSoft's recommended guidelines, how should the System API data model be defined?

- ☐ A. The System APIs should expose all back-end system fields
- ☒ B. The data model of the System APIs should make use of data types that approximately mirror those from the back-end systems
- ☐ C. If there are misspellings of the data fields in the back-end system, System APIs should not correct it, and expose it as-is to mirror the back-end systems
- ☐ D. The data model should define its own naming convention, and not follow the same naming as the back-end systems

Question 14 of 60



Choose 1 option.

A customer wants to monitor and gain insights about the number of requests coming in a given time period as well as to measure key performance indicators (response times, CPU utilization, number of active APIs).

Which tool provides these data insights?

- ☒ A. Anypoint Monitoring
- ☐ B. Runtime Alerts
- ☐ C. Functional Monitoring
- ☐ D. API Manager

Question 15 of 60



Choose 1 option.

An organization has built an application network following the API-led connectivity approach recommended by MuleSoft. To protect the application network against attacks from malicious external API clients, the organization plans to apply JSON Threat Protection policies.

To which API-led connectivity layer should the JSON Threat Protection policies most commonly be applied?

- ☐ A. Process layer
- ☒ B. Experience layer
- ☐ C. All layers
- ☐ D. System layer

Question 16 of 60



Choose 3 options.

Which three tools automate the deployment of Mule applications?

- ☐ A. API Community Manager
- ☐ B. Anypoint Studio
- ☒ C. Platform APIs
- ☒ D. Anypoint Platform CLI
- ☐ E. Runtime Manager
- ☒ F. Mule Maven plugin

Question 17 of 60



Choose 2 options.

A Production environment is running on a dedicated Virtual Private Cloud (VPC) on CloudHub 1.0, and the security team guidelines clearly state no traffic on HTTP.

Which two options support these security guidelines?

- ☒ A. Remove the entry from the VPC firewall rule

```
{  
  "CIDR Block": "0.0.0.0/0", // (Anywhere)  
  "Protocol": "TCP",  
  "From port": 8081,  
},  
  
{  
  "CIDR Block": "10.111.0.0/24", // (Local VPC)  
  "Protocol": "TCP",  
  "From port": 8091,  
}
```

- ☐ B. Configure the IP Blocklist policy to control access to a configured API endpoint from either a single IP address or a range of IP addresses

Question 18 of 60



Choose 1 option.

An organization requires several APIs to be secured with OAuth 2.0, and PingFederate has been identified as the identity provider for API client authorization. The PingFederate Client Provider is configured in access management, and the PingFederate OAuth 2.0 Token Enforcement policy is configured for the API instances required by the organization. The API instances reside in two business groups (Group A and Group B) within the Master Organization (Master Org).

What should be done to allow API consumers to access the API instances?

- ☒ A. The API consumer should create a client application and request access to the API in Anypoint Exchange, and the API administrator should approve the request
- ☐ B. The API consumer should create a client application and request access to the API in Ping Identity, and the organization's Ping Identity workflow will grant access
- ☐ C. The API administrator should grant access to the API consumers by creating contracts in the relevant API instances in API Manager
- ☐ D. The API administrator should configure the correct client discovery URL in both child business groups, and the API consumer should request access to the API in Ping Identity

Question 19 of 60



Choose 1 option.

Which statement is true about Spike Control policy and Rate Limiting policy?

- ☒ A. All requests are rejected after the limit is reached in Rate Limiting policy, whereas the requests are queued in Spike Control policy after the limit is reached
- ☐ B. To protect Experience APIs by limiting resource consumption, Rate Limiting policy must be applied
- ☐ C. In a clustered environment, the Rate Limiting and Spike Control policies are applied to each node in the cluster
- ☐ D. In order to apply Rate Limiting and Spike Control policies, a contract to bind client application and API is needed for both

Question 20 of 60



Choose 1 option.

A large lending company has developed an API to unlock data from a database server and web server. The API has been deployed to Anypoint Virtual Private Cloud (VPC) on CloudHub 1.0.

The database server and web server are in the customer's secure network and are not accessible through the public internet. The database server is in the customer's AWS VPC, whereas the web server is in the customer's on-premises corporate data center.

How can access be enabled for the API to connect with the database server and the web server?

- ☐ A. Set up VPC peering with the customer's on-premises corporate data center and a VPN tunnel to AWS VPC
- ☒ B. Set up VPC peering with AWS VPC and a VPN tunnel to the customer's on-premises corporate data center
- ☐ C. Set up VPC peering with AWS VPC and the customer's on-premises corporate data center
- ☐ D. Set up a transit gateway to the customer's on-premises corporate data center through AWS VPC

Question 21 of 60



Choose 1 option.

Which statement is true about identity management and client management on Anypoint Platform?

- ☐ A. If an external identity provider is configured, the SAML 2.0 bearer tokens issued by the identity provider cannot be used for invocations of the Anypoint Platform web APIs
- ☒ B. If an external client provider is configured, it must be configured at the Anypoint Platform organization level and cannot be assigned to individual business groups and environments
- ☐ C. Anypoint Platform supports configuring one external identity provider
- ☐ D. Both client management and identity management require an identity provider

Question 22 of 60



Choose 1 option.

An online store's marketing team has noticed an increase in customers leaving online baskets without checking out. They suspect a technology issue is at the root cause of the baskets being left behind. They approach the Center for Enablement to ask for help identifying the issue. Multiple APIs from across all the layers of their application network are involved in the shopping application.

Which feature of the Anypoint Platform can be used to view metrics from all involved APIs at the same time?


- ☐ A. API Manager
- ☐ B. Custom dashboards
- ☒ C. Built-in dashboards
- ☐ D. Functional monitoring

Question 23 of 60



Choose 1 option.

Refer to the exhibits. Which architectural constraint is compatible with the API-led connectivity architectural style?



- ☒ A. Use a Process API to orchestrate calls to multiple System APIs but not to other Process APIs



Question 24 of 60



Choose 1 option.

A customer wants to host their MuleSoft applications in CloudHub 1.0, and these applications should be available at the domain `https://api.acmecorp.com`.

After creating a dedicated load balancer (DLB) called `acme-dlb-prod`, which further action must the customer take to complete the configuration?

- ☐ A. Configure the DLB with a TLS certificate for `api.acmecorp.com` and create a CNAME record from `api.acmecorp.com` to `acme-dlb-prod.e1.cloudbhub.io`
- ☐ B. Configure the DLB with a TLS certificate for `api.acmecorp.com` and create an A record for `api.acmecorp.com` to the public IP addresses associated with their DLB
- ☐ C. Configure the DLB with a TLS certificate for `acme-dlb-prod.lb.anypointdns.net` and create a CNAME record from `api.acmecorp.com` to `acme-dlb-prod.lb.anypointdns.net`
- ☒ D. Configure the DLB with a TLS certificate for `api.acmecorp.com` and create a CNAME record from `api.acmecorp.com` to `acme-dlb-prod.lb.anypointdns.net`

Question 25 of 60




Choose 1 option.

Refer to the exhibits.

The Line of Business (LoB) of an eCommerce company is requesting a process that sends automated notifications via email every time a new order is processed through the customer's mobile application or through the internal company's web application. In the future, multiple notification channels may be added: for example, text messages and push notifications.

What is the most effective API-led connectivity approach for the scenario described above?

- ☒ A. Create one Experience API for the web application and one for the mobile application. Create a Process API to orchestrate, retrieve the email template from a database, and send the email using the Anypoint Connector for Email.



Question 26 of 60



Choose 1 option.

A Platform Architect inherits a legacy monolithic SOAP-based web service that performs a number of tasks, including showing all policies belonging to a client. The service connects to two back-end systems – a life-insurance administration system and a general-insurance administration system – and then queries for insurance policy information within each system, aggregates the results, and presents a SOAP-based response to a user interface (UI).

The architect wants to break up the monolithic web service to follow API-led conventions.

Which part of the service should be put into the process layer?

- ☒ A. Combining the insurance policy information from the administration systems
- ☐ B. Authenticating and maintaining connections to each of the back-end administration systems
- ☐ C. Querying the data from the administration systems
- ☐ D. Presenting the SOAP-based response to the UI

Question 27 of 60



Choose 1 option.

An API with multiple API implementations (Mule applications) is deployed to both CloudHub and customer-hosted Mule runtimes. All the deployments are managed by the MuleSoft-hosted control plane. An alert needs to be triggered whenever an API implementation stops responding to API requests, even if no API clients have called the API implementation for some time.

What is the most effective out-of-the-box solution to create these alerts to monitor the API implementations?

- ☒ A. Create monitors in Anypoint Functional Monitoring for the API implementations, where each monitor repeatedly invokes an API implementation endpoint
- ☐ B. Handle API invocation exceptions within the calling API client and raise an alert from that API client when such an exception is thrown
- ☐ C. Add code to each API client to send an Anypoint Platform REST API request to generate a custom alert in Anypoint Platform when an API invocation times out
- ☐ D. Configure one Worker Not Responding alert in Anypoint Runtime Manager for all API implementations that will then monitor every API implementation

Question 28 of 60



Choose 1 option.

An eCommerce company is adding a new Product Details feature to their website. A customer will launch the product catalog page, a new Product Details link will appear by product where they can click to retrieve the product detail description. Product detail data is updated with product update releases, once or twice a year. Presently the database response time has been very slow due to high volume.

What action retrieves the product details with the lowest response time, fault tolerant, and consistent data?

- ☒ A. Select the product details from a database in a Cache scope and return them within the API response
- ☐ B. Select the product details from a database and return them within the API response
- ☐ C. Select the product details from a database and put them in Anypoint MQ; the Anypoint MQ subscriber will receive the product details and return them within the API response
- ☐ D. Use an object store to store and retrieve the product details originally read from a database and return them within the API response

Question 29 of 60



Choose 1 option.

Which component monitors APIs and endpoints at scheduled intervals, receives reports about whether tests pass or fail, and displays statistics about API and endpoint performance?

- ☒ A. API Functional Monitoring
- ☐ B. API Analytics
- ☐ C. Anypoint Monitoring dashboards
- ☐ D. Anypoint Runtime Manager alerts

Question 30 of 60



Choose 1 option.

Which out-of-the-box key performance indicator measures the success of a typical Center for Enablement and is immediately available in responses from Anypoint Platform APIs?

- ☐ A. Per deployed API implementation, the amount of bandwidth consumed each day
- ☒ B. Per published API, the number of consumers that requested access to the API and have been approved in the Production environment
- ☐ C. Per business group, the ratio of the number of production API implementations deployed using a CI/CD pipeline to the number of production API implementations deployed manually
- ☐ D. Per published API, the number of developers that downloaded a version of the API specification

Question 31 of 60



Choose 1 option.

Which scenario is suited for MUnit tests instead of integration tests?

- ☒ A. For read-only interactions to any dependencies (such as other web APIs)
- ☐ B. When no mocking is permissible
- ☐ C. For tests that are implemented using SoapUI
- ☐ D. When testing does not require knowledge of implementation details

Question 32 of 60



Choose 1 option.

A large organization with an experienced central IT department is getting started using MuleSoft. There is a project to connect a siloed back-end system to a new Customer Relationship Management (CRM) system. The Center for Enablement is coaching them to use API-led connectivity.

What action would support the creation of an application network using API-led connectivity?

- ☐ A. To expedite this project, central IT should extend the CRM system and back-end systems to connect to one another using built in integration interfaces
- ☐ B. Determine if the new CRM system supports the creation of custom REST APIs, establishes a private network with CloudHub, and supports OAuth 2.0 authentication
- ☐ C. Invite the business analyst to create a business process model to specify the canonical data model between the two systems
- ☒ D. Create a System API to unlock the data on the back-end system using a REST API

Question 33 of 60



Choose 1 option.

A large company wants to implement IT infrastructure in its own data center, based on the corporate IT policy requirements that data and metadata reside locally.

Which combination of Mule control plane and Mule runtime plane(s) meets the requirements?

- ☐ A. The MuleSoft-hosted control plane and Anypoint Runtime Fabric for the runtime plane
- ☒ B. Anypoint Platform Private Cloud Edition for the control plane and customer-hosted Mule runtimes for the runtime plane
- ☐ C. The MuleSoft-hosted control plane and customer-hosted Mule runtimes for the runtime plane
- ☐ D. Anypoint Platform Private Cloud Edition for the control plane and the MuleSoft-hosted runtime plane


Question 34 of 60



Choose 1 option.

Refer to the exhibit.

Which APIs can be used with DataGraph to create a unified schema?



Question 35 of 60



Choose 1 option.

When should idempotency be taken into account?

- ☐ A. When storing the results of a previous request for use in response to subsequent requests
- ☒ B. When preventing duplicate processing from multiple sent requests
- ☐ C. When sending concurrent update requests for the same entity
- ☐ D. When making requests to update currently locked entities

Question 36 of 60



Choose 1 option.

An API is protected with a Client ID Enforcement policy and uses the default configuration. Access is requested for the client application to the API, and an approved contract now exists between the client application and the API

How can a consumer of this API avoid a 401 error "Unauthorized or invalid client application credentials"?

- ☒ A. Send the obtained token as a header in every call
- ☐ B. Send the obtained *client_id* and *client_secret* in the request body
- ☐ C. Send the obtained *client_id* and *client_secret* in the header of every API Request call
- ☐ D. Send the obtained *client_id* and *client_secret* as URI parameters in every call

Question 37 of 60



Choose 2 options.

Which two statements are true about the technology architecture of an Anypoint Virtual Private Cloud (VPC)?

- ☒ A. CIDR blocks are used
- ☐ B. Round-robin load balancing is used to distribute client requests across different applications
- ☐ C. By default, HTTP requests can be made from the public internet to workers at port 8091
- ☒ D. Ports 8081 and 8082 are used
- ☐ E. Anypoint VPC is responsible for load balancing the applications

Question 38 of 60



Choose 1 option.

A company deploys Mule applications with default configurations through Runtime Manager to customer-hosted Mule runtimes. Each Mule application is an API implementation that exposes RESTful interfaces to API clients. The Mule runtimes are managed by the MuleSoft-hosted control plane. The payload is never used by any Logger components.

When an API client sends an HTTP request to a customer-hosted Mule application, which metadata or data (payload) is pushed to the MuleSoft-hosted control plane?

- ☒ A. Only the metadata
- ☐ B. The data and metadata
- ☐ C. No data
- ☐ D. Only the data

Question 39 of 60



Choose 1 option.

A company is using an on-prem cluster in the data center as a runtime plane and MuleSoft-hosted control plane.

How can the company monitor the detailed performance metrics on the Mule applications deployed to the cluster from the control plane?

- ☐ A. Due to the potential performance impact on the runtime nodes, the Monitoring agent should be installed on a separate server
- ☒ B. There is no action needed as the on-prem runtime automatically sends the performance data to the control plane
- ☐ C. Monitoring Agent must be installed on each node in the cluster
- ☐ D. The settings of the Monitoring section in the control plane must be updated to enable detailed logging on the metrics to be captured

Question 40 of 60




Choose 1 option.

Refer to the exhibits.

A company stores financial transaction data in two legacy systems. For each legacy system, a separate, dedicated System API (SAPI) exposes data for that legacy system. A Process API (PAPI) merges the data retrieved from all of the System APIs into a common format. Several API clients call the PAPI through its public domain name.

The company now wants to expose a subset of financial data to a newly developed mobile application that uses a different Bounded Context Data Model. The company wants to follow MuleSoft's best practices for building out an effective application network.

Following MuleSoft's best practices, how can the company expose financial data needed by the mobile application in a way that minimizes the impact on the currently running API clients, API implementations, and support asset reuse?




Question 41 of 60



Choose 1 option.

Refer to the exhibit.

Which API policy would least likely be used when designing an Experience API that is intended to work with a consumer mobile phone or tablet application?



- ☐ A. Client ID Enforcement
- ☐ B. JSON Threat Protection
- ☒ C. IP Allowlist
- ☐ D. OAuth 2.0 Access Token Enforcement

Question 40 of 60




Choose 1 option.

Refer to the exhibits.

A company stores financial transaction data in two legacy systems. For each legacy system, a separate, dedicated System API (SAPI) exposes data for that legacy system. A Process API (PAPI) merges the data retrieved from all of the System APIs into a common format. Several API clients call the PAPI through its public domain name.

The company now wants to expose a subset of financial data to a newly developed mobile application that uses a different Bounded Context Data Model. The company wants to follow MuleSoft's best practices for building out an effective application network.

Following MuleSoft's best practices, how can the company expose financial data needed by the mobile application in a way that minimizes the impact on the currently running API clients, API implementations, and support asset reuse?



Question 40 of 60



Choose 1 option.


Legacy System-1

Legacy System-2

- ☐ D. Refactor the PAPI implementation into an Experience API (EAPI).

Add data transformation and routing logic to the EAPI implementation to support the required endpoints of both mobile and web clients.

Postpone deployment until all API consumers acknowledge they are ready to migrate to the new API version.




Question 41 of 60



Choose 1 option.

Refer to the exhibit.

Which API policy would least likely be used when designing an Experience API that is intended to work with a consumer mobile phone or tablet application?



- ☐ A. Client ID Enforcement
- ☐ B. JSON Threat Protection
- ☒ C. IP Allowlist
- ☐ D. OAuth 2.0 Access Token Enforcement

Question 42 of 60



Choose 1 option.

Refer to the exhibit.

A Mule 4 API has been deployed to CloudHub and a Basic Authentication - Simple policy has been applied to all API methods and resources. However, the API is still accessible by clients without using authentication.

How is this possible?

Apply Basic authentication - Simple policy

Enforces HTTP Basic authentication according to the details configured in the policy

User Name *

User Password *

Method & Resource conditions

- ☒ Apply configurations to all API methods & resources
- ☐ Apply configurations to specific methods & resources

Cancel

Apply

Question 43 of 60



Choose 1 option.

A Platinum customer uses the U.S. control plane and deploys applications to CloudHub in Singapore with a default log configuration.

The compliance officer asks where the logs and monitoring data reside?

- ☐ A. Logs and monitoring data are held in the United States
- ☐ B. Logs and monitoring data are held in Singapore
- ☐ C. Logs are held in the United States and monitoring data is held in Singapore
- ☒ D. Logs are held in Singapore and monitoring data is held in the United States



Question 42 of 60



Choose 1 option.

Refer to the exhibit.

A Mule 4 API has been deployed to CloudHub and a Basic Authentication - Simple policy has been applied to all API methods and resources. However, the API is still accessible by clients without using authentication.

How is this possible?

Apply Basic authentication - Simple policy

Enforces HTTP Basic authentication according to the details configured in the policy

User Name *

User Password *

Method & Resource conditions

- ☒ Apply configurations to all API methods & resources
- ☐ Apply configurations to specific methods & resources

Cancel

Apply

Question 42 of 60



Choose 1 option.

User Password *

Method & Resource conditions

- ☒ Apply configurations to all API methods & resources
- ☐ Apply configurations to specific methods & resources

Cancel

Apply

- ☐ A. The API Router component is pointing to the incorrect Exchange version of the API
- ☒ B. The Autodiscovery element is not present in the deployed Mule application
- ☐ C. No contracts for client applications have been created for this API
- ☐ D. One of the application's CloudHub workers restarted

Question 43 of 60



Choose 1 option.

A Platinum customer uses the U.S. control plane and deploys applications to CloudHub in Singapore with a default log configuration.

The compliance officer asks where the logs and monitoring data reside?

- ☐ A. Logs and monitoring data are held in the United States
- ☐ B. Logs and monitoring data are held in Singapore
- ☐ C. Logs are held in the United States and monitoring data is held in Singapore
- ☒ D. Logs are held in Singapore and monitoring data is held in the United States

Question 44 of 60



Choose 1 option.

8012.

should be on the agenda of the first meeting?

- ☒ A. Define C4E objectives, mission statement, guiding principles, and responsibilities
- ☐ B. Specify operating model for the MuleSoft Integrations division
- ☐ C. A walk through of common-services best practices for logging, auditing, exception handling, caching, security via policy, and rate limiting/throttling via policy
- ☐ D. Explore API monetization options based on identified use cases through MuleSoft

Question 45 of 60



Choose 1 option.

A company is building an application network using MuleSoft's recommendations for various API layers.

What is the main (default) role of a process API in an application network?

- ☐ A. To manage and process the secure direct communication between a back-end system and an end-user client or mobile device in the application network
- ☐ B. To secure and optimize the data synchronization processing of large data dumps between back-end systems
- ☒ C. To automate parts of business processes by coordinating and orchestrating the invocation of other APIs in the application network
- ☐ D. To secure, manage, and process communication with specific types of end-user client applications or devices in the application network

Question 46 of 60




Choose 1 option.

Refer to the exhibit.

A new upstream API is being designed to offer a service-level agreement (SLA) of 500ms median and 800ms maximum (99th percentile) response time. The corresponding API implementation must sequentially invoke three downstream APIs of equivalent complexity.

The first of these downstream APIs offers the following SLA for its response time. Median: 100ms; 80th percentile: 500ms; 95th percentile: 1,000ms.


If possible, how can a timeout be set in the upstream API for the invocation of the first downstream API to meet the new upstream API's desired SLA; if this cannot be done, what should be done to best support the new upstream API?



Question 46 of 60



Choose 1 option.



- ☐ A. Set a timeout of 100ms; this leaves 400ms for the other two downstream APIs to complete
- ☐ B. Set a timeout of 50ms; this times out more invocations of that API but gives additional room for retries
- ☐ C. Do not set a timeout; the invocation of this API is mandatory and requires waiting until it responds
- ☒ D. No timeout is possible to meet the upstream API's desired SLA; a different SLA must be negotiated with the first downstream API or an alternative API must be invoked

Question 47 of 60



Choose 1 option.

An auto manufacturer has a mature CI/CD practice and wants to automate packaging and deployment of any Mule applications to various deployment targets, including CloudHub workers/replicas, customer-hosted Mule runtimes, and Anypoint Runtime Fabric.

Which MuleSoft-provided tool or component facilitates automating the packaging and deployment of Mule applications to various deployment targets as part of the company's CI/CD practice?

- ☐ A. Anypoint Runtime Manager
- ☐ B. Anypoint Platform CLI
- ☐ C. Anypoint Platform REST APIs
- ☒ D. Mule Maven plugin

Question 48 of 60



Choose 1 option.

Several times a week, an API implementation shows several thousand requests per minute in an Anypoint Monitoring dashboard. Between these bursts, the dashboard shows between two and five requests per minute. The API implementation is running on Anypoint Runtime Fabric with two non-clustered replicas, reserved vCPU 1.0 and vCPU Limit 2.0.

An API consumer has complained about slow response time, and the dashboard shows the 99 percentile is greater than 120 seconds at the time of the complaint. It also shows greater than 90% CPU usage during these time periods.

In manual tests in the QA environment, the API consumer has consistently reproduced the slow response time and high CPU usage, and there were no other API requests at this time. In a brainstorming session, the engineering team has created several proposals to reduce the response time for requests.

Which proposal should be pursued first?

- ☐ A. Modify the API client to split the problematic request into smaller, less-demanding requests
- ☐ B. Throttle the API client to reduce the number of requests per minute
- ☒ C. Increase the number of replicas of the API implementation
- ☐ D. Increase the vCPU resources of the API implementation

Question 49 of 60



Choose 1 option.

A Rate Limiting policy is applied to an API implementation to protect the back-end system. Recently, there have been surges in demand that cause some API client POST requests to the API implementation to be rejected with policy-related errors, causing delays and complications to the API clients.

How should the API policies that are applied to the API implementation be changed to reduce the frequency of errors returned to API clients, while still protecting the back-end system?

- ☐ A. Keep the Rate Limiting policy and add an SLA-based Spike Control policy
- ☒ B. Remove the Rate Limiting policy and add a Spike Control policy
- ☐ C. Keep the Rate Limiting policy and add a Client ID Enforcement policy
- ☐ D. Remove the Rate Limiting policy and add an HTTP Caching policy

Question 50 of 60



Choose 2 options.

An existing Quoting API is defined in RAML and used by REST clients for interacting with the quoting engine. Currently there is a resource defined in the RAML that allows the creation of quotes; however, a new requirement was just received to allow for the updating of existing quotes.

Which two actions need to be taken to facilitate this change so it can be processed?

- ☐ A. Deprecate existing versions of the API in Exchange
- ☒ B. Update the RAML with new method details for the update request
- ☐ C. Add a new API policy to API Manager to allow access to the updated endpoint
- ☐ D. Remove the old client applications and create new client applications to account for the changes
- ☒ E. Update the API implementation to accommodate the new update request

Question 51 of 60



Choose 2 options.

An IT Security Compliance Auditor is assessing which nonfunctional requirements (NFRs) are already being implemented to meet security measures.

- The Web API has Rate-Limiting SLA
- Basic Authentication - LDAP
- JSON Threat Protection
- IP Allowlist policies applied

Which two NFRs are enforced?

- ☐ A. The API is protected against XML invocation attacks
- ☐ B. Sensitive data is masked to prevent compromising critical information
- ☒ C. The API invocations are coming from a known subnet range
- ☒ D. Username/password supported to validate login credentials
- ☐ E. Performance expectations are to be allowed up to 1,000 requests per second

Question 52 of 60



Choose 1 option.

A business process is being implemented within an organization's application network. The architecture group proposes using a more coarse-grained application network design with relatively fewer APIs deployed to the application network compared to a more fine-grained design.

Overall, which factor typically increases with a more coarse-grained design for this business process implementation and deployment compared with using a more fine-grained design?

- ☒ A. The complexity of each API implementation
- ☐ B. The usage of network infrastructure resources by the application network
- ☐ C. The number of discoverable assets related to APIs deployed in the application network
- ☐ D. The number of possible connections between API implementations in the application network

Question 53 of 60




Choose 1 option.

Refer to the exhibit.

An Order API triggers a sequence of other API calls to look up details of an order's items in a back-end inventory database. The Order API calls the OrderItems process API, which calls the Inventory system API. The Inventory system API performs database operations in the back-end inventory database.

The network connection between the Inventory system API and the database is known to be unreliable and hang at unpredictable times.

Where should a two-second timeout be configured in the API processing sequence so that the Order API never waits more than two seconds for a response from the OrderItems process API?



- ☐ A. In the OrderItems process API implementation
- ☒ B. In the Inventory system API implementation

Question 54 of 60



Choose 1 option.

A TemperatureSensors API instance is defined in API Manager in the PROD environment of the CAR_FACTORY business group. An AcmeTemperatureSensors Mule application implements this API instance and is deployed from Runtime Manager to the PROD environment of the CAR_FACTORY business group. A policy that requires a valid client ID and client secret is applied in API Manager to the API instance.

Where can an API consumer obtain a valid client ID and client secret to call the AcmeTemperatureSensors Mule application?

- ☐ A. In API Manager, from the PROD environment of the CAR_FACTORY business group
- ☐ B. In access management, from the PROD environment of the CAR_FACTORY business group
- ☐ C. In Anypoint Exchange, from an API client application that has been approved for the TemperatureSensors API instance
- ☒ D. In secrets manager, request access to the Shared Secret static username/password

Question 53 of 60




Choose 1 option.

Refer to the exhibit.

An Order API triggers a sequence of other API calls to look up details of an order's items in a back-end inventory database. The Order API calls the OrderItems process API, which calls the Inventory system API. The Inventory system API performs database operations in the back-end inventory database.

The network connection between the Inventory system API and the database is known to be unreliable and hang at unpredictable times.

Where should a two-second timeout be configured in the API processing sequence so that the Order API never waits more than two seconds for a response from the OrderItems process API?



- ☐ A. In the OrderItems process API implementation
- ☒ B. In the Inventory system API implementation

Question 56 of 60



Choose 1 option.

A European company has customers all across Europe, and the IT department is migrating from an older platform to MuleSoft. The main requirements are that the new platform should allow redeployments with zero downtime and deployment of applications to multiple runtime versions, provide security and speed, and utilize Anypoint MQ as the message service.

Which runtime plane should the company select based on the requirements without additional network configuration?

- ☐ A. Anypoint Runtime Fabric on Self-Managed Kubernetes for the runtime plane
- ☐ B. Customer-hosted runtime plane
- ☒ C. MuleSoft-hosted runtime plane (CloudHub)
- ☐ D. Runtime Fabric on VMs / Bare Metal for the runtime plane

Question 55 of 60



Choose 1 option.

The asset version 2.0.0 of the Order API is successfully published in Exchange and configured in API Manager with the Autodiscovery API ID correctly linked to the API implementation. A new GET method is added to the existing API specification, and after updates, the asset version of the Order API is 2.0.1.

What happens to the Autodiscovery API ID when the new asset version is updated in API Manager?

- ☐ A. The API ID changes, so the API implementation must be updated with the latest API ID for the new asset version in the API Autodiscovery global element
- ☐ B. The API ID does not change, but the API implementation must be updated in the API Autodiscovery global element to indicate the new asset version 2.0.1
- ☐ C. The API ID changes, but no changes are needed to the API implementation for the new asset version in the API Autodiscovery global element because the API ID is automatically updated
- ☒ D. The API ID does not change, so no changes to the API implementation are needed for the new asset version in the API Autodiscovery global element

Question 57 of 60



Choose 1 option.

A team is planning to enhance an Experience API specification, and they are following API-led connectivity design principles.

What is their motivation for enhancing the API?

- ☐ A. The underlying System API is updated to provide more detailed data for several heavily used resources
- ☐ B. An IP Allowlist policy is being added to the API instances in the Development and Staging environments
- ☐ C. A Canonical Data Model is being adopted that impacts several types of data included in the API
- ☒ D. The primary API consumer wants certain kinds of endpoints changed from the Center for Enablement standard to the consumer system standard

Question 58 of 60



Choose 1 option.

A client has several applications running on the Salesforce service cloud. The business requirement for integration is to get daily data changes from Account and Case Objects. Data needs to be moved to the client's private cloud AWS DynamoDB instance as a single JSON and the business foresees only wanting five attributes from the Account object, which has 219 attributes (some custom) and eight attributes from the Case Object.

What design should be used to support the API/ Application data model?

- ☐ A. Create separate entities for Account and Case Objects by mimicking all the attributes in SAPI, which are combined by the PAPI and filtered to provide JSON output containing 13 attributes.
- ☐ B. Start implementing an Enterprise Data Model by defining enterprise Account and Case Objects and implement SAPI and DynamoDB tables based on the Enterprise Data Model.
- ☐ C. Request client's AWS project team to replicate all the attributes and create Account and Case JSON table in DynamoDB. Then create separate entities for Account and Case Objects by mimicking all the attributes in SAPI to transfer JSON data to DynamoDB for respective Objects
- ☒ D. Create separate entities for Account with five attributes and Case with eight attributes in SAPI, which are combined by the PAPI to provide JSON output containing 13 attributes.

Question 59 of 60



Choose 1 option.

A manufacturing company has deployed an API implementation to CloudHub and has not configured it to be automatically restarted by CloudHub when the worker is not responding.

Which statement is true when no API Client invokes that API implementation?

- ☐ A. Alerts on the API invocations are raised but no alerts on the API implementation can be raised
- ☐ B. Alerts on the API invocation and API implementation can be raised
- ☒ C. No alert on the API invocations is raised but alerts on the API implementation can be raised
- ☐ D. No alert on the API invocations and API implementation can be raised

Question 60 of 60



Choose 1 option.

An operations team is analyzing the effort needed to set up monitoring of their application network. They are looking at which API invocation metrics can be used to identify and predict trouble without having to write custom scripts or install additional analytics software or tools.

Which type of metrics can satisfy this goal of directly identifying and predicting failures?

- ☐ A. The effectiveness of the application network based on the level of reuse
- ☒ B. The number and types of API policy violations per day
- ☐ C. The number and types of past API invocations across the application network
- ☐ D. The ROI from each API invocation